

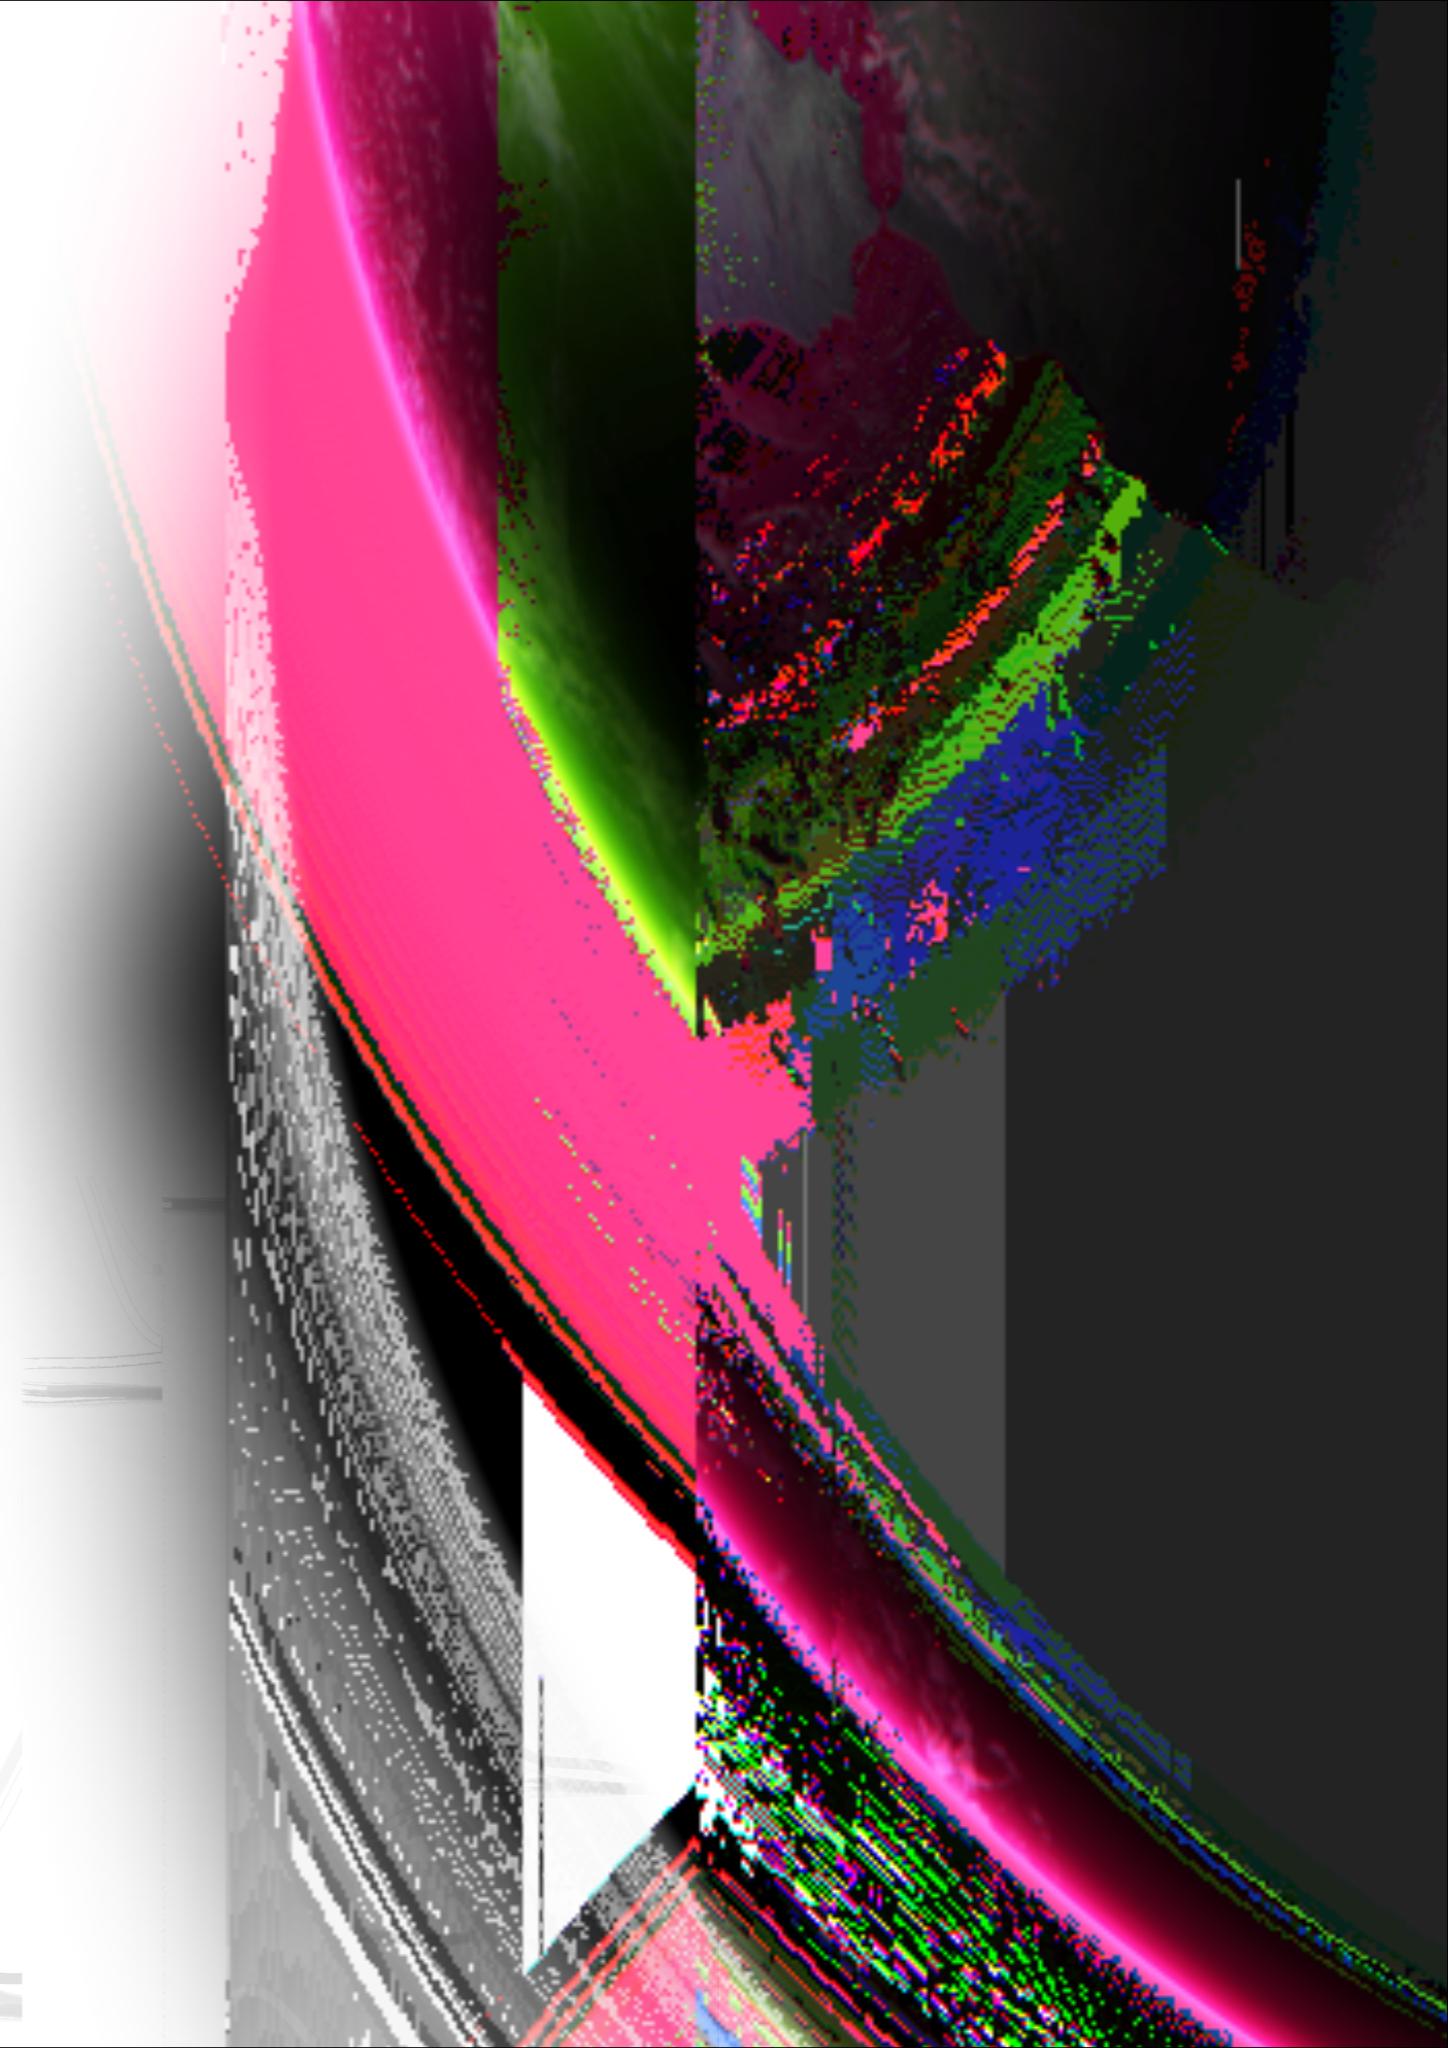
Teachin

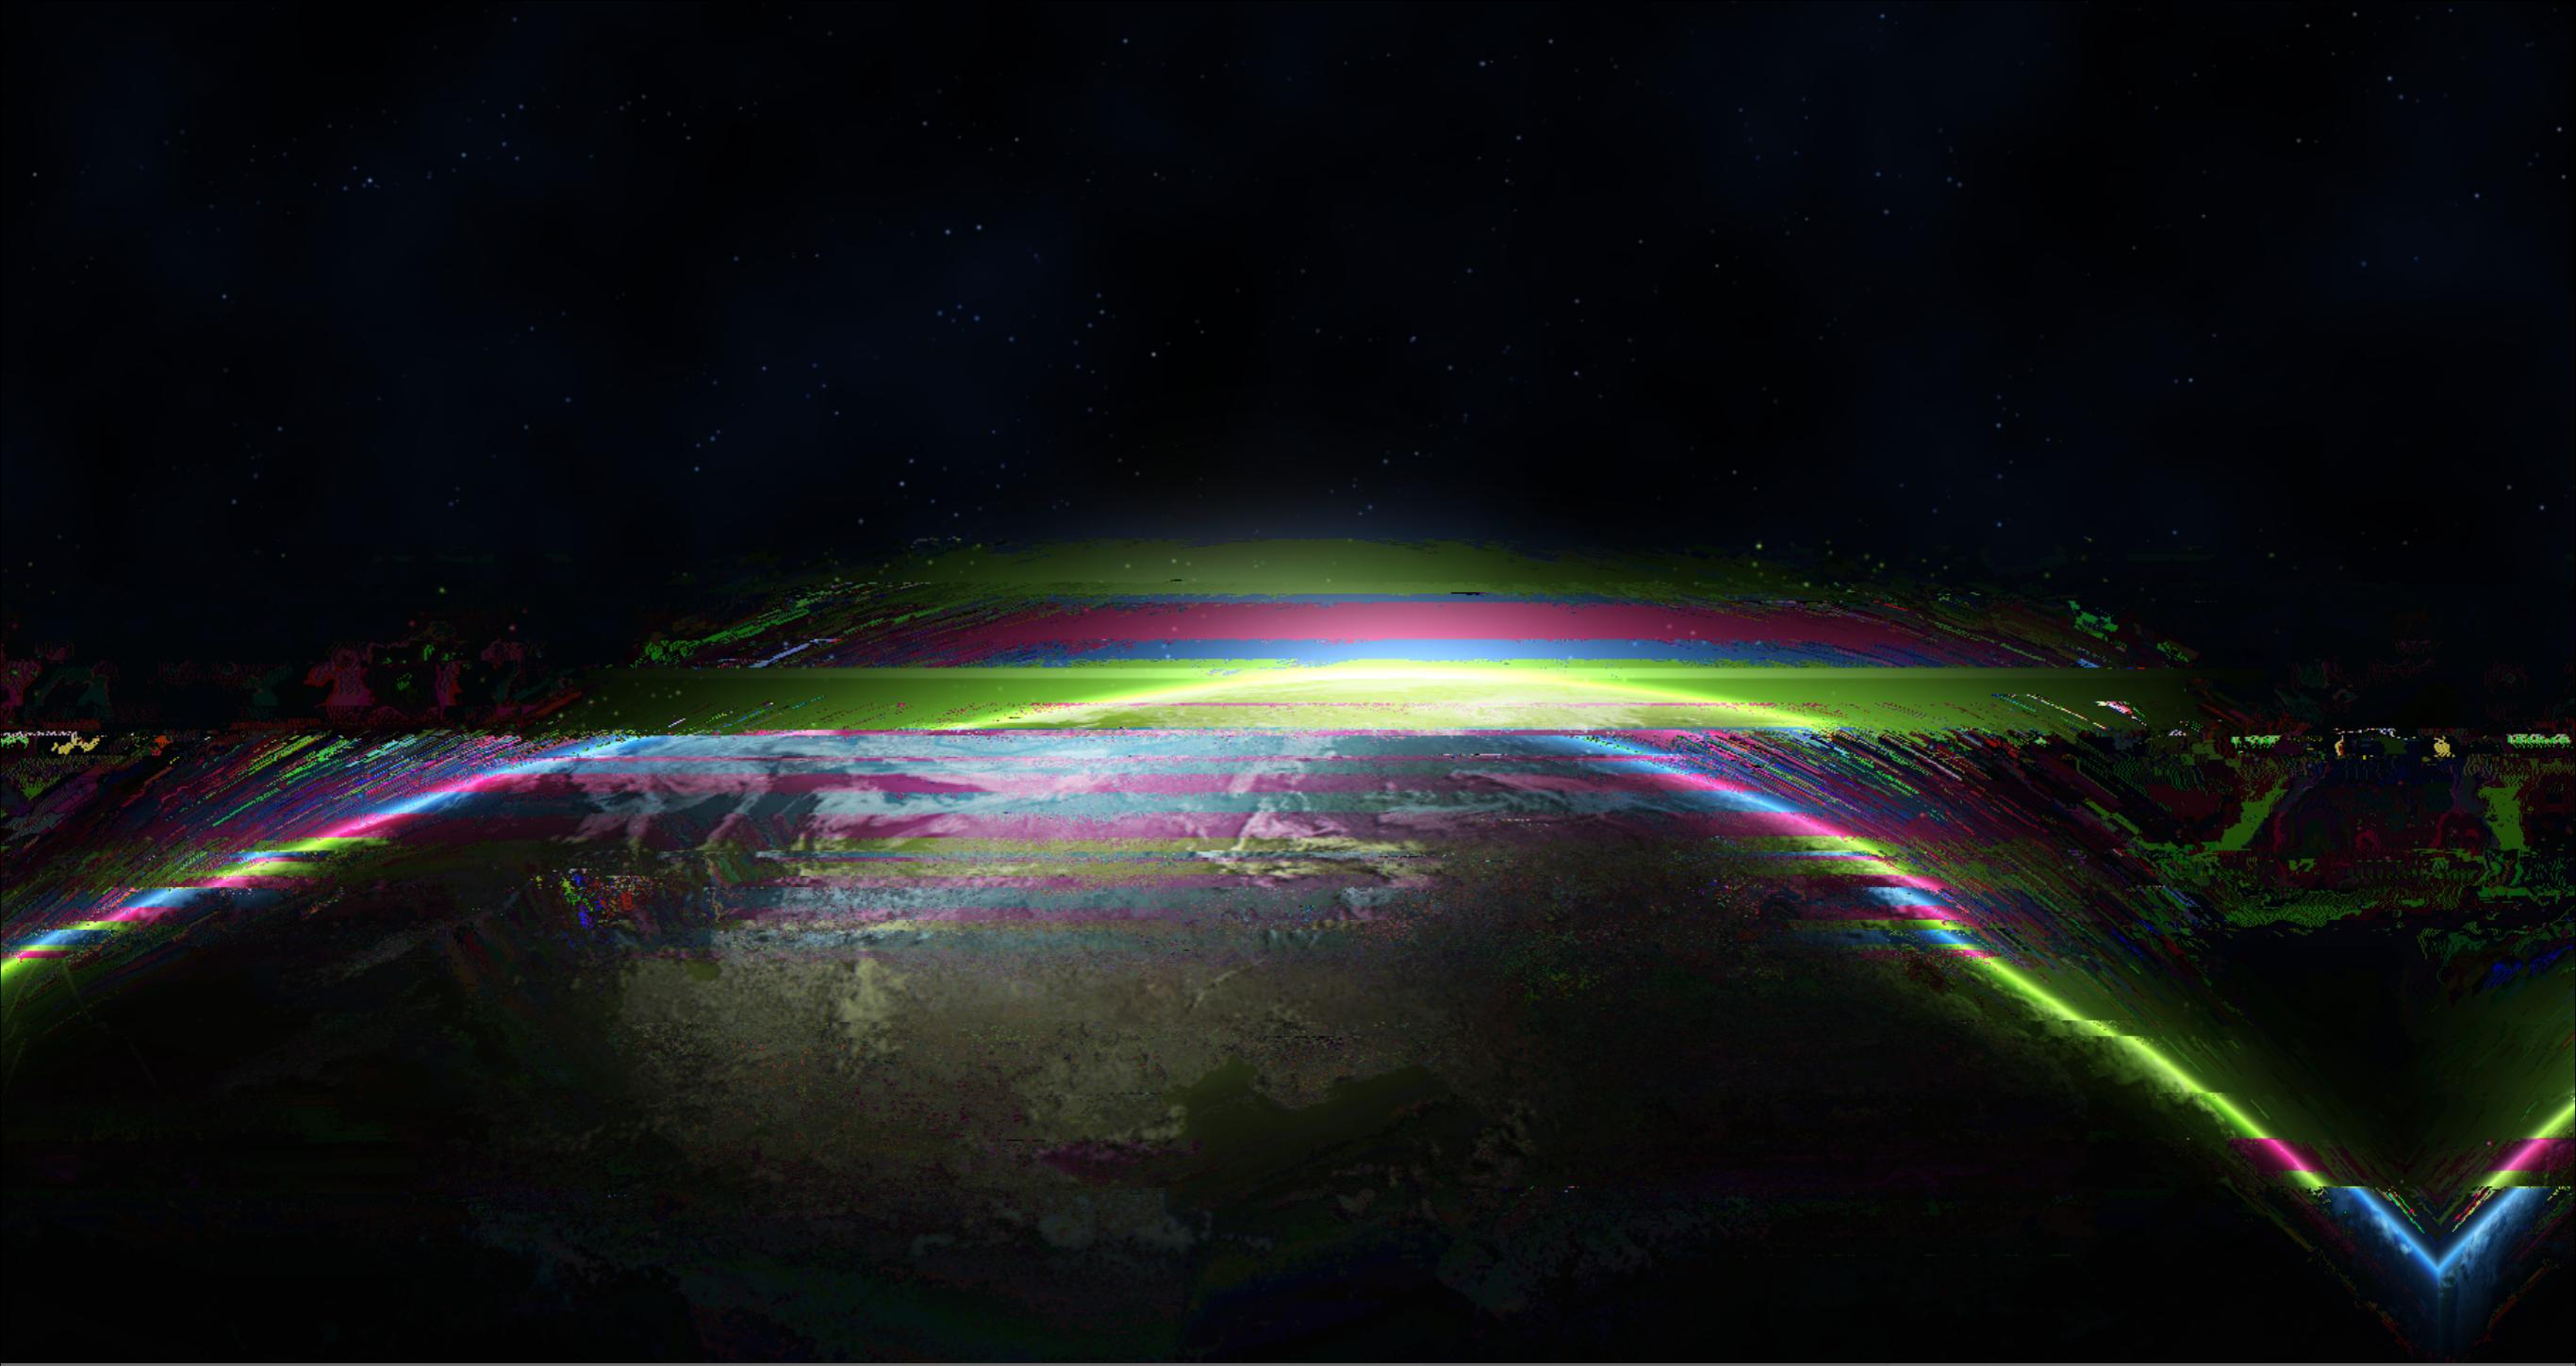
Abd Da a

CE408 at Scale

Overview

- ❖ IMPACT
- ❖ Hi
- ❖ Achieve
- ❖ Sedata
- ❖ Le
- ❖ Fe

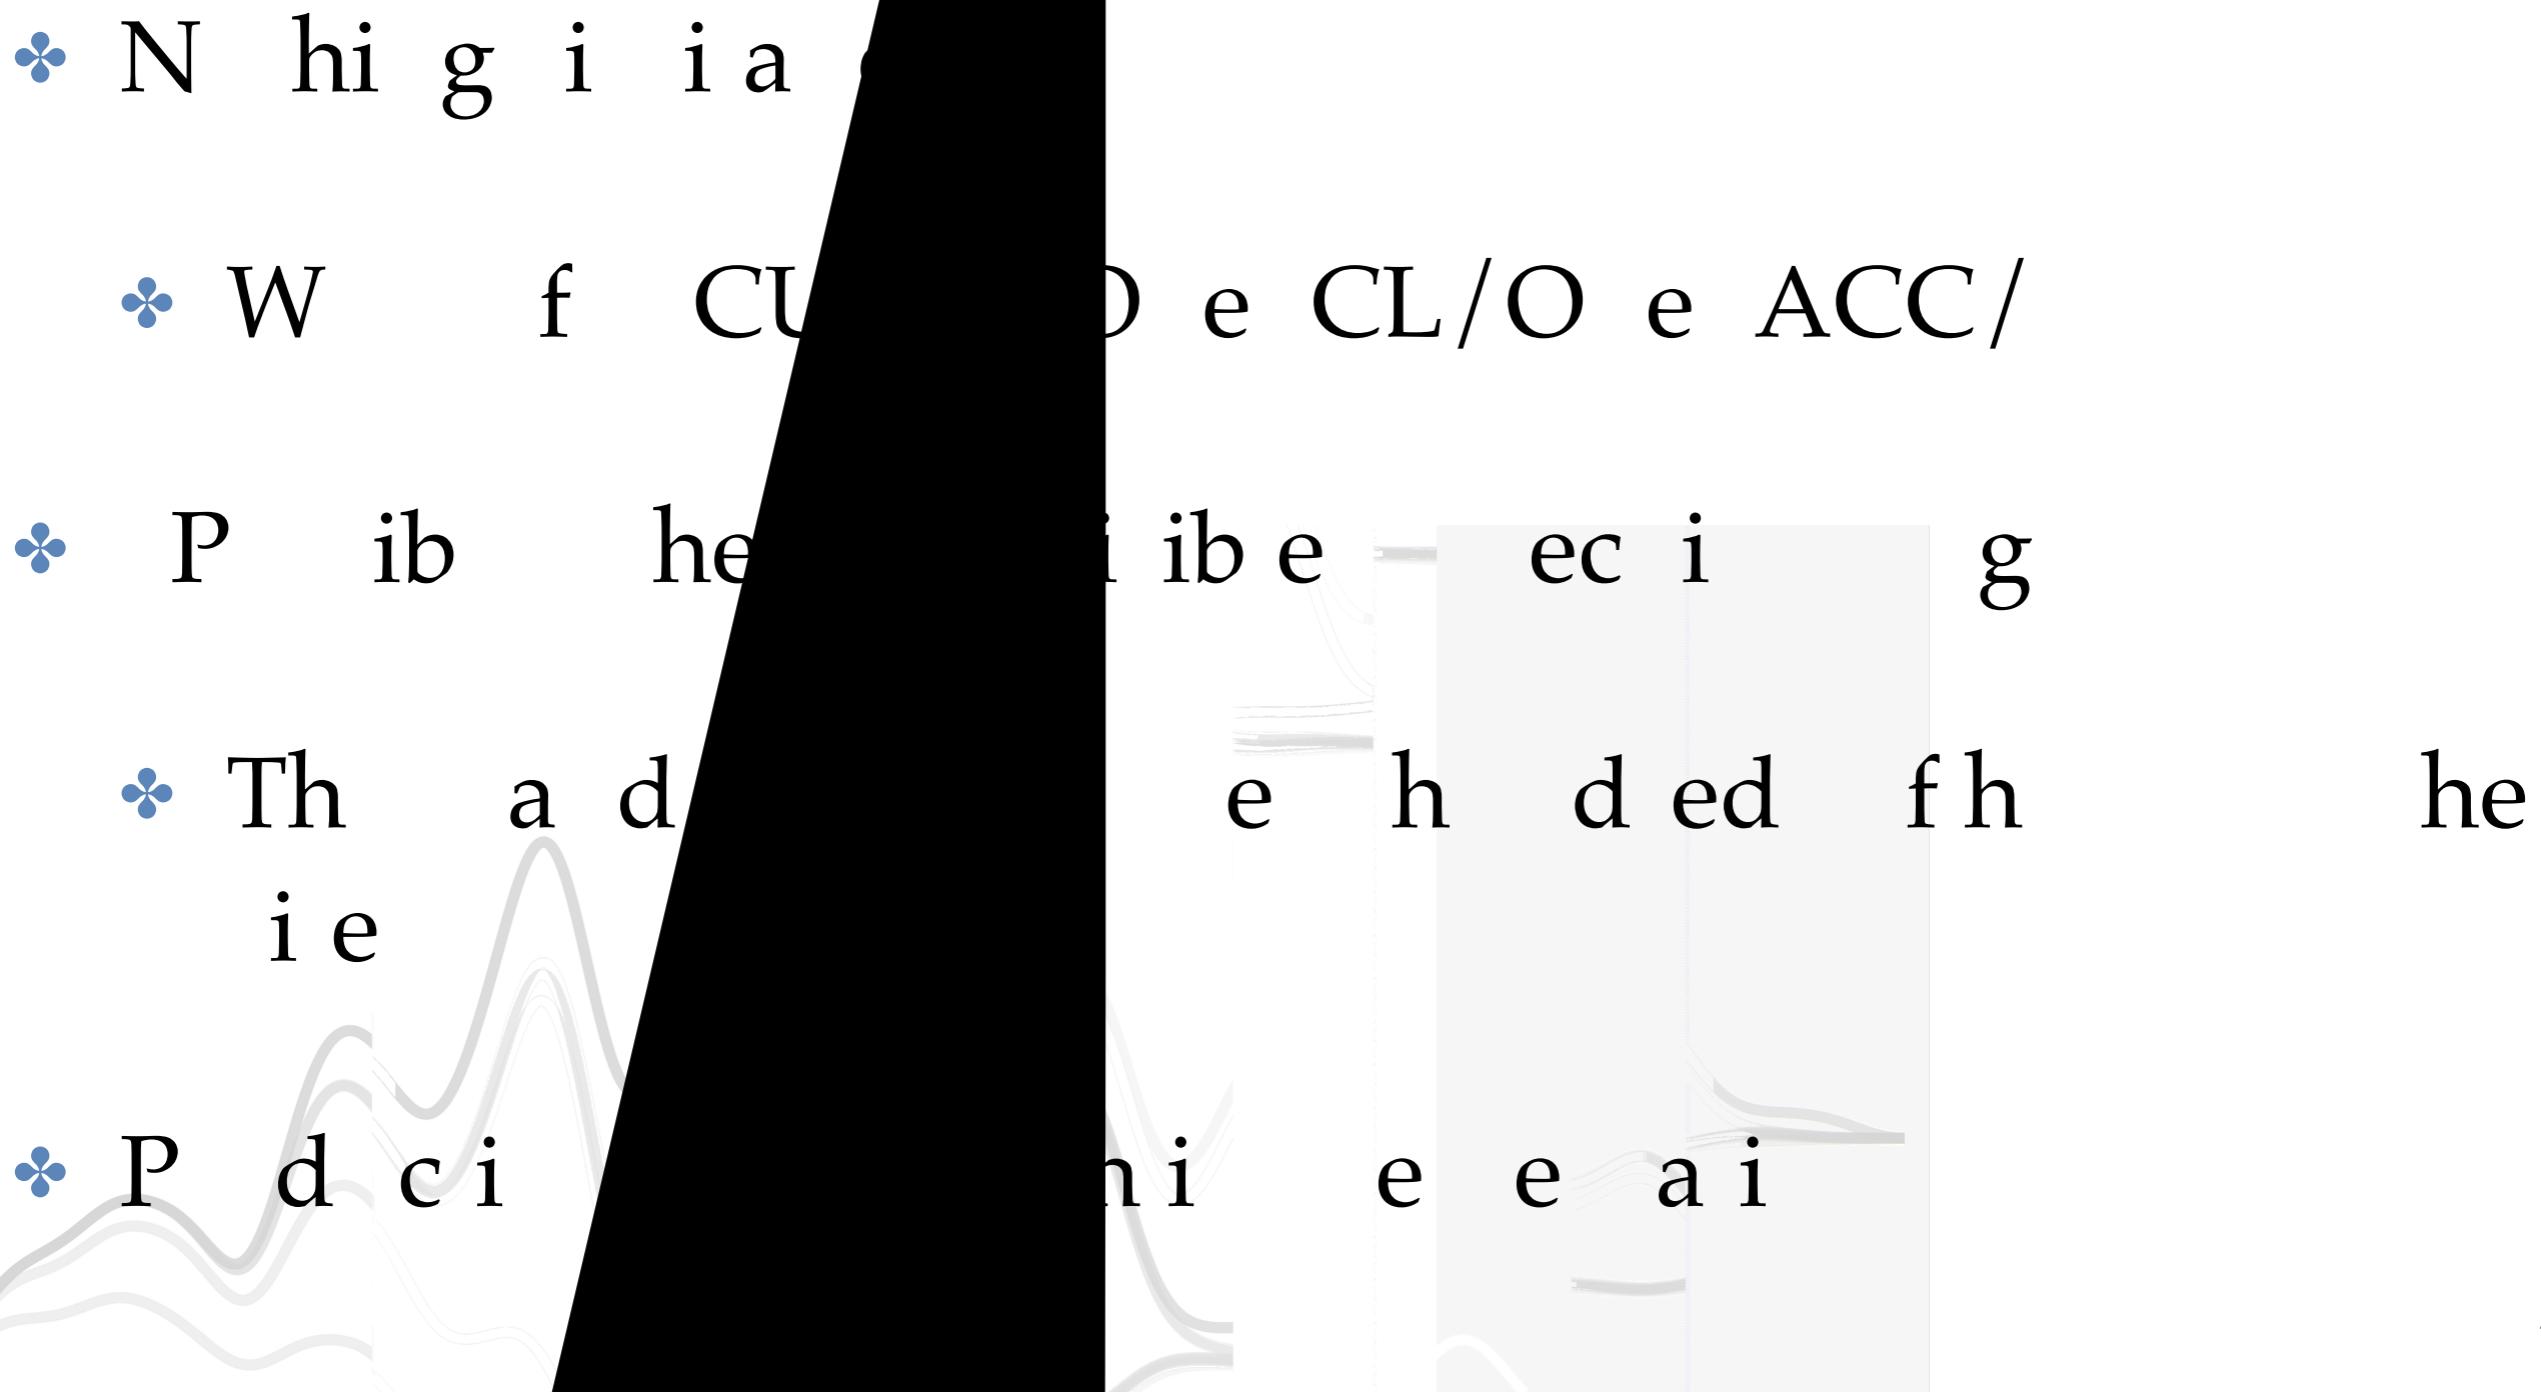




Quick Demo

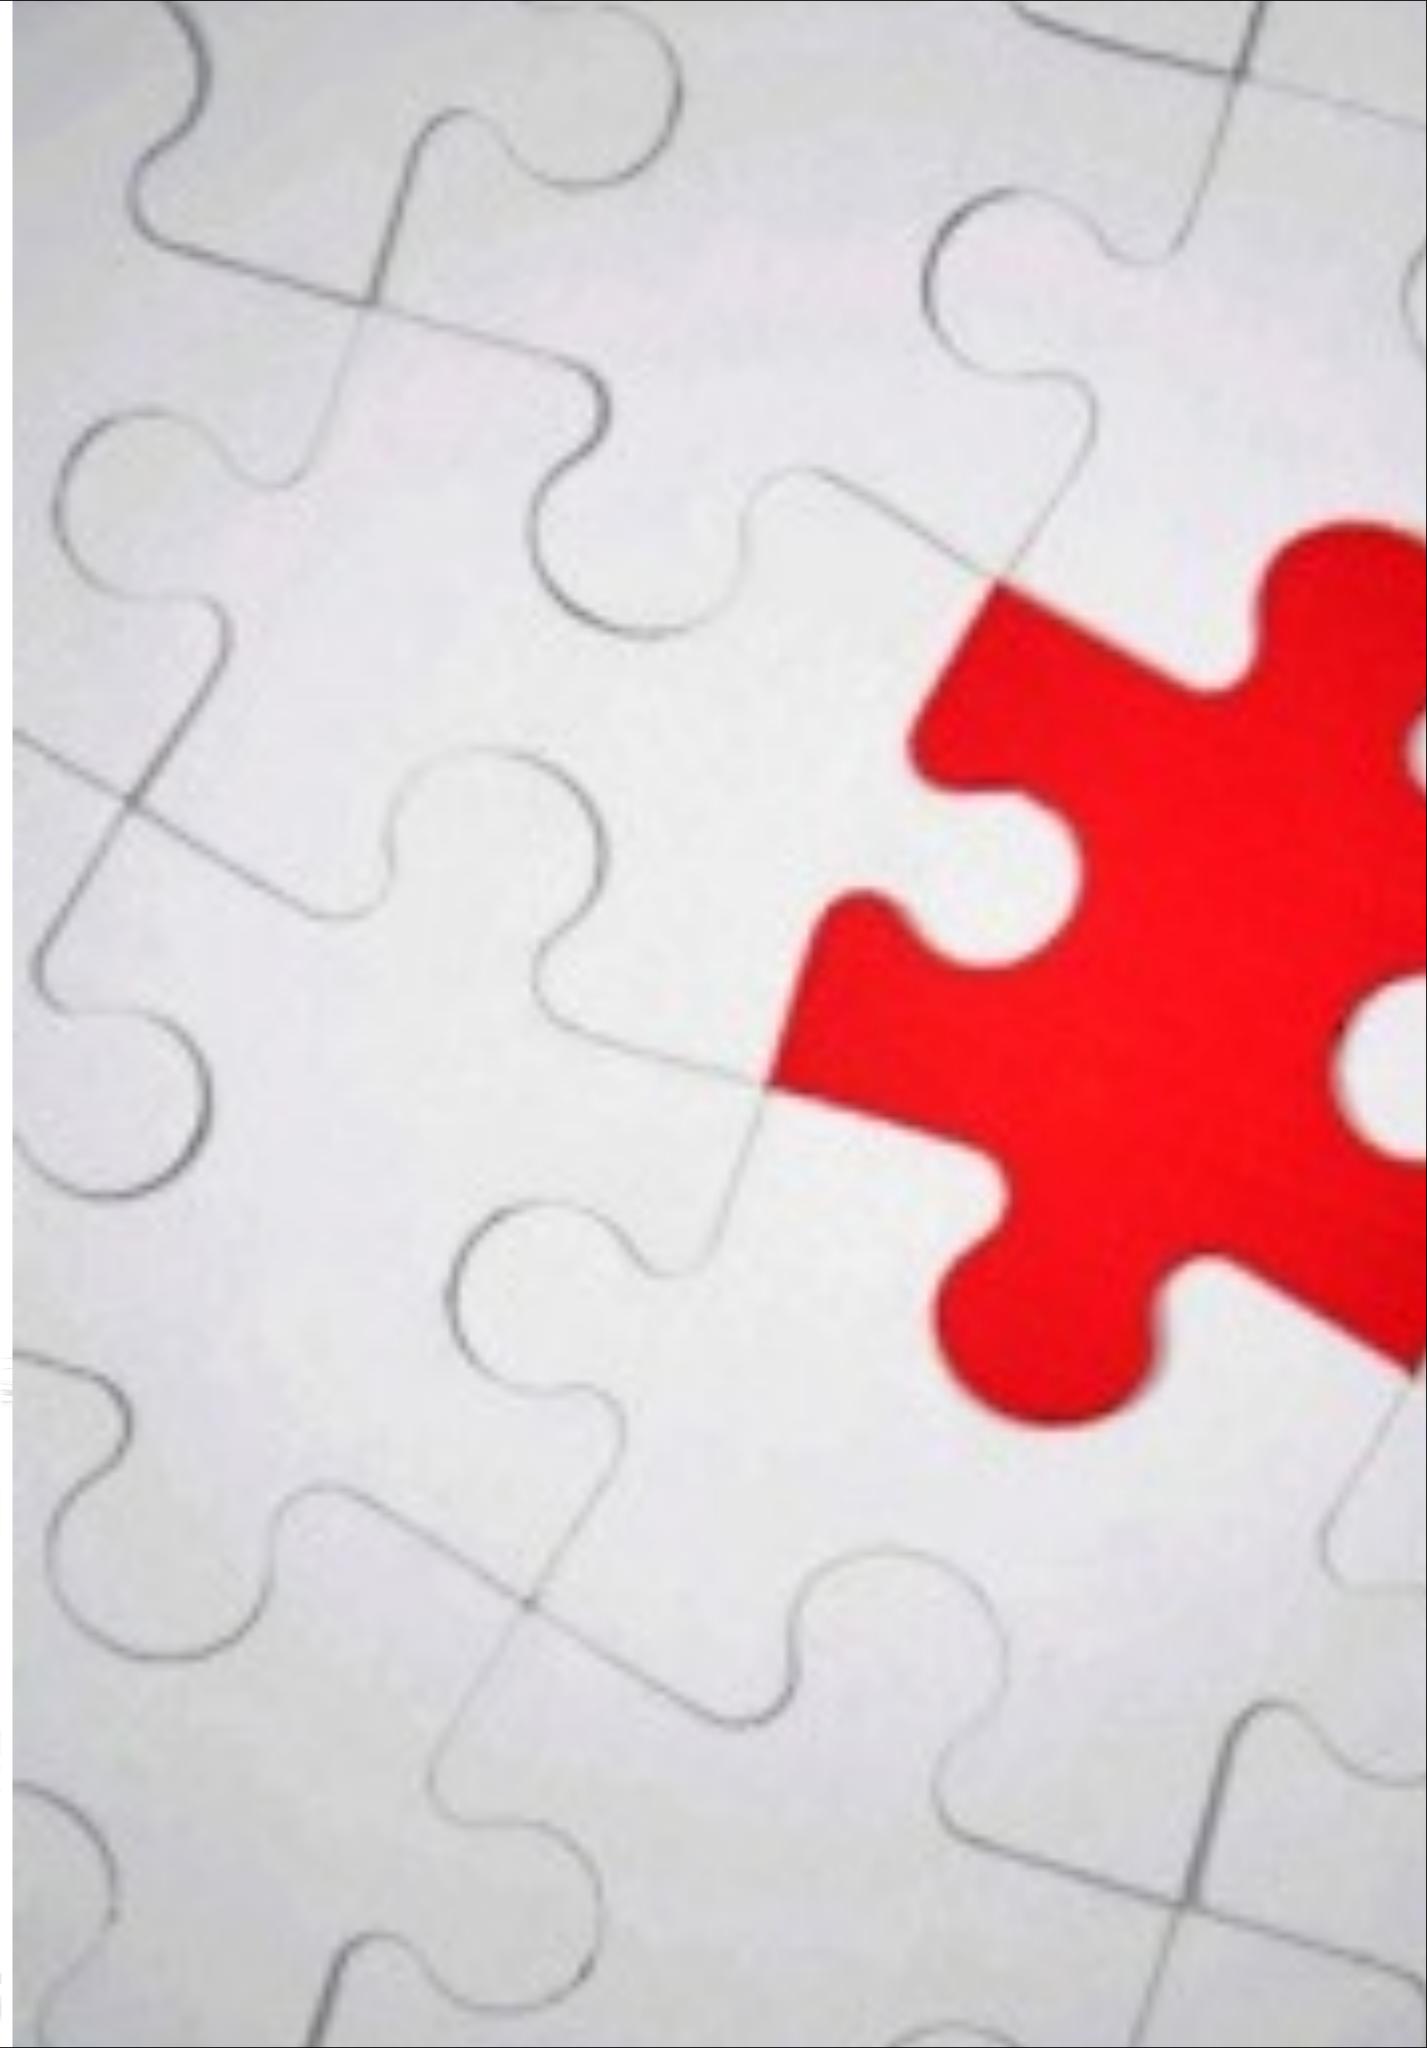
IMPACT

The
IMPACT
Research Group



Objective

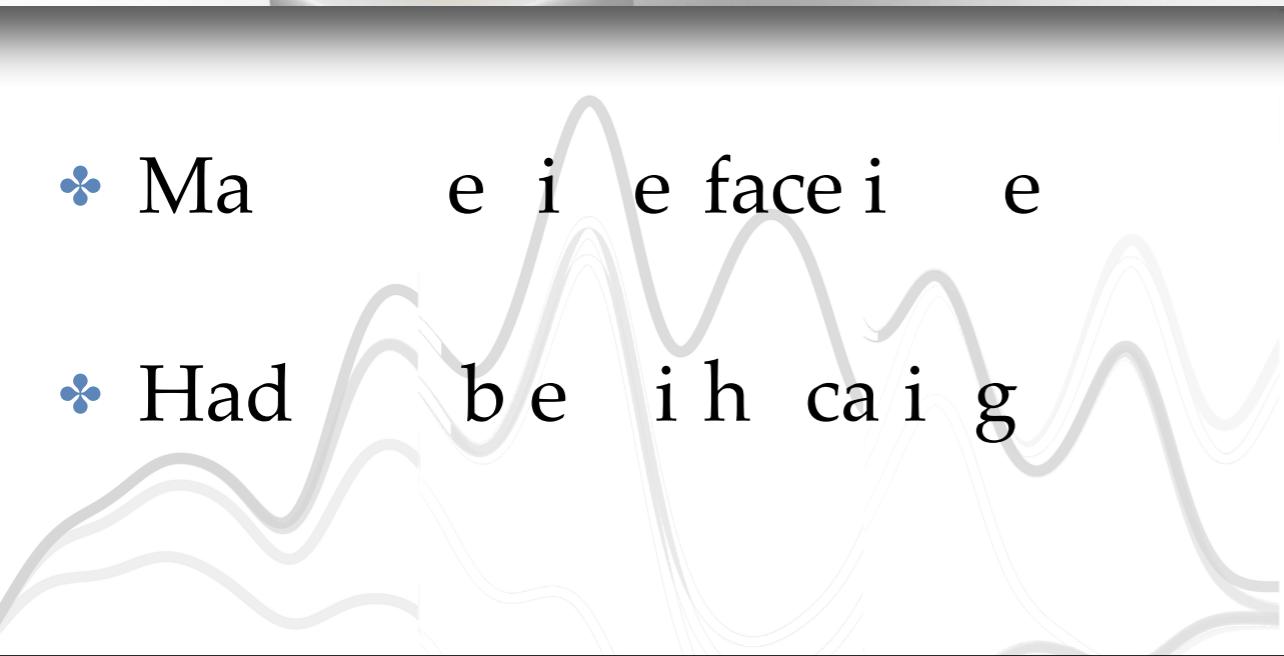
- ❖ Create a game idea C
e i e f he C e a
c e
- ❖ A e e de e
ide f he e i e
- ❖ N be ied i C e a
(a ffe i f e
ch)



Previous System



- ✿ Ma e i e face i e
- ✿ Had b e i h c a i g



- ✿ G adi g a d e f i e
- ✿ Did ha e ee e ie i g

New System

- ✿ I a g adi g and
b i i bac
- ✿ Pee e ie i
- ✿ Sca i g a
- ✿ Ma i g a e i

C

e

i



How do you grade 4000 Programs?

Y d .

Architecture

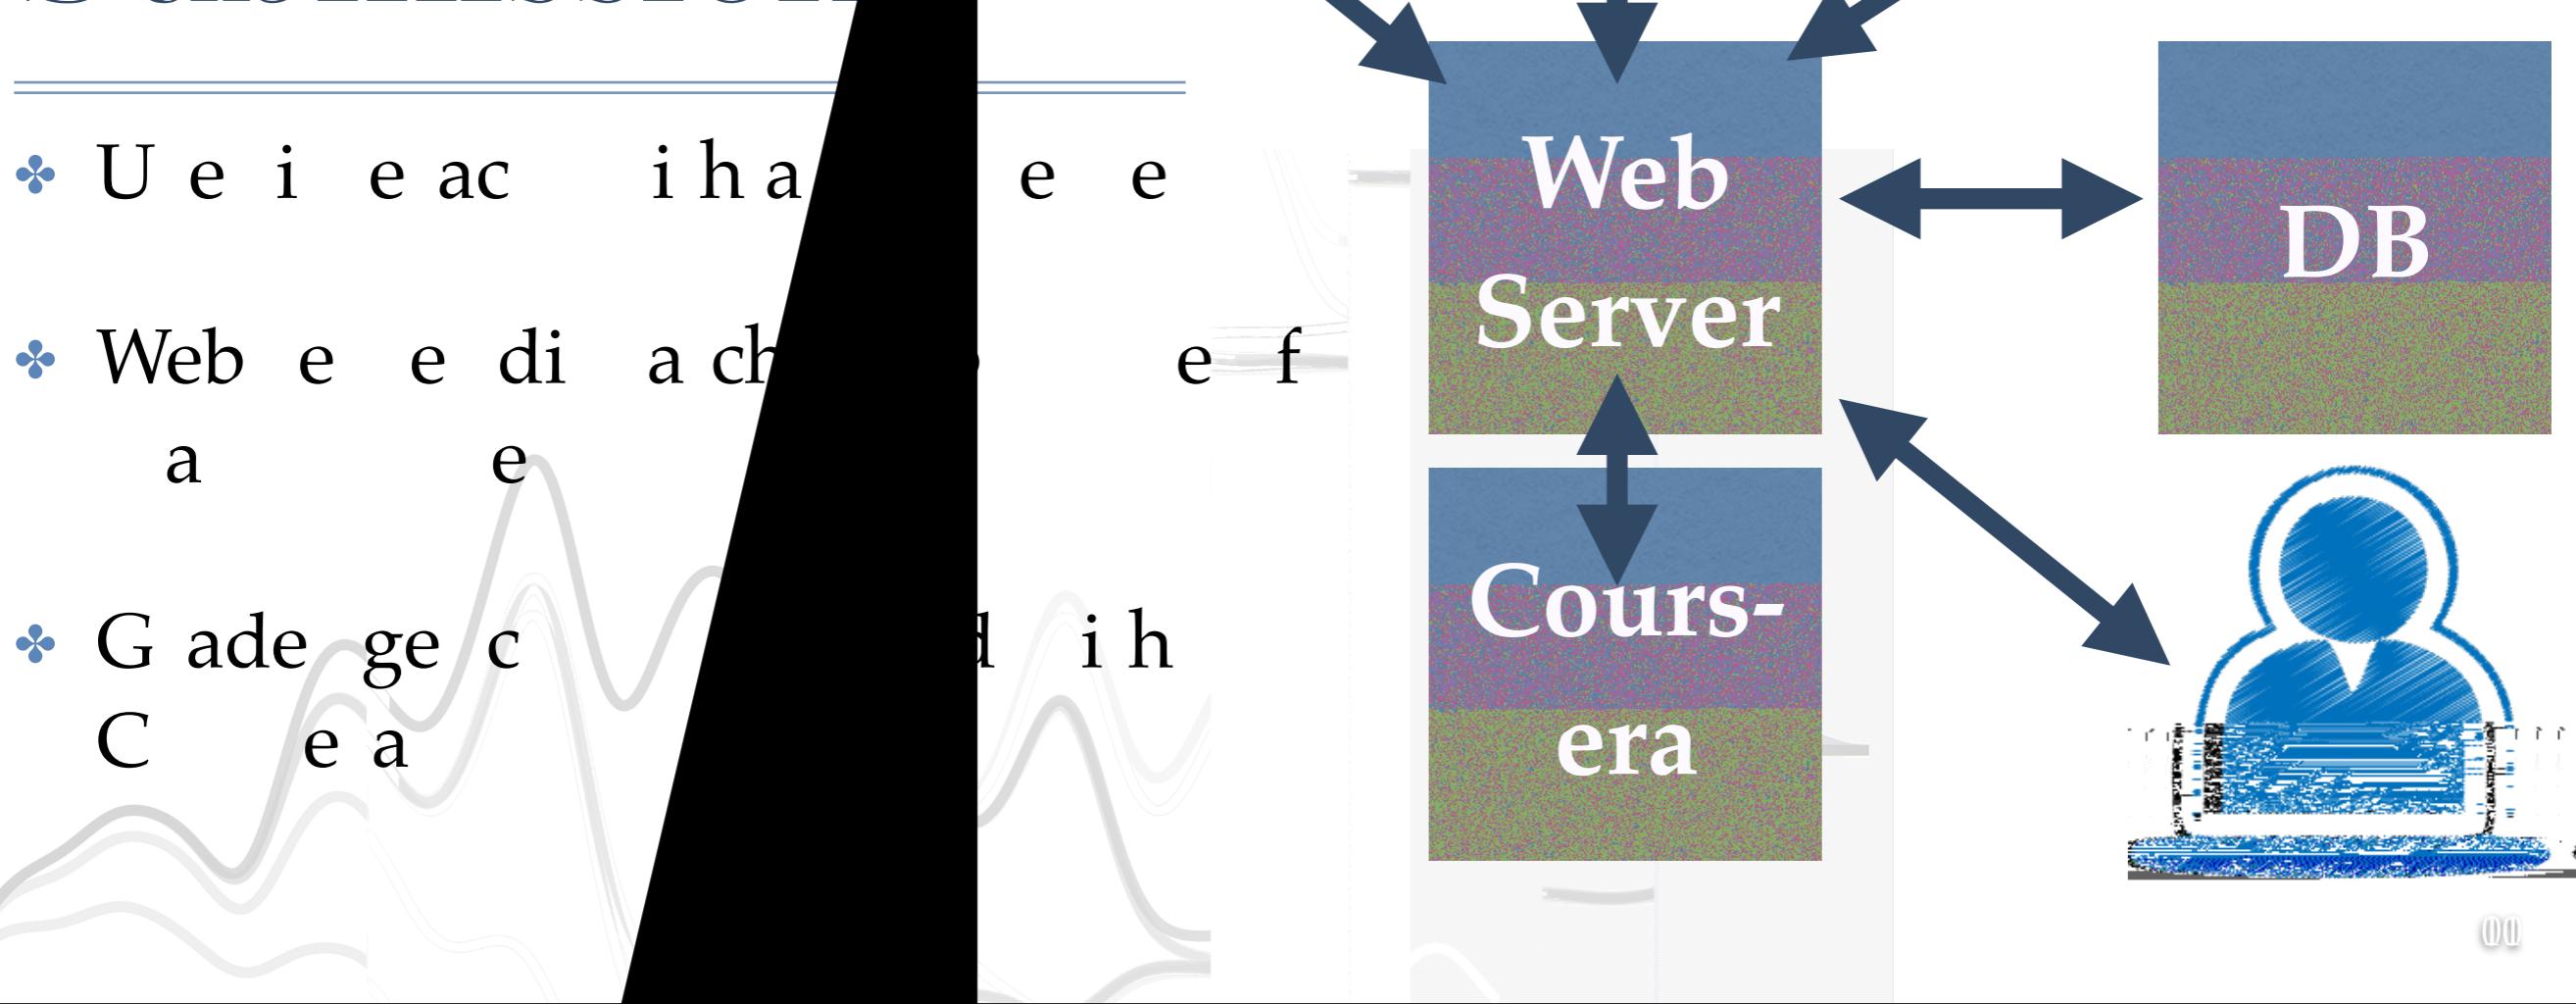


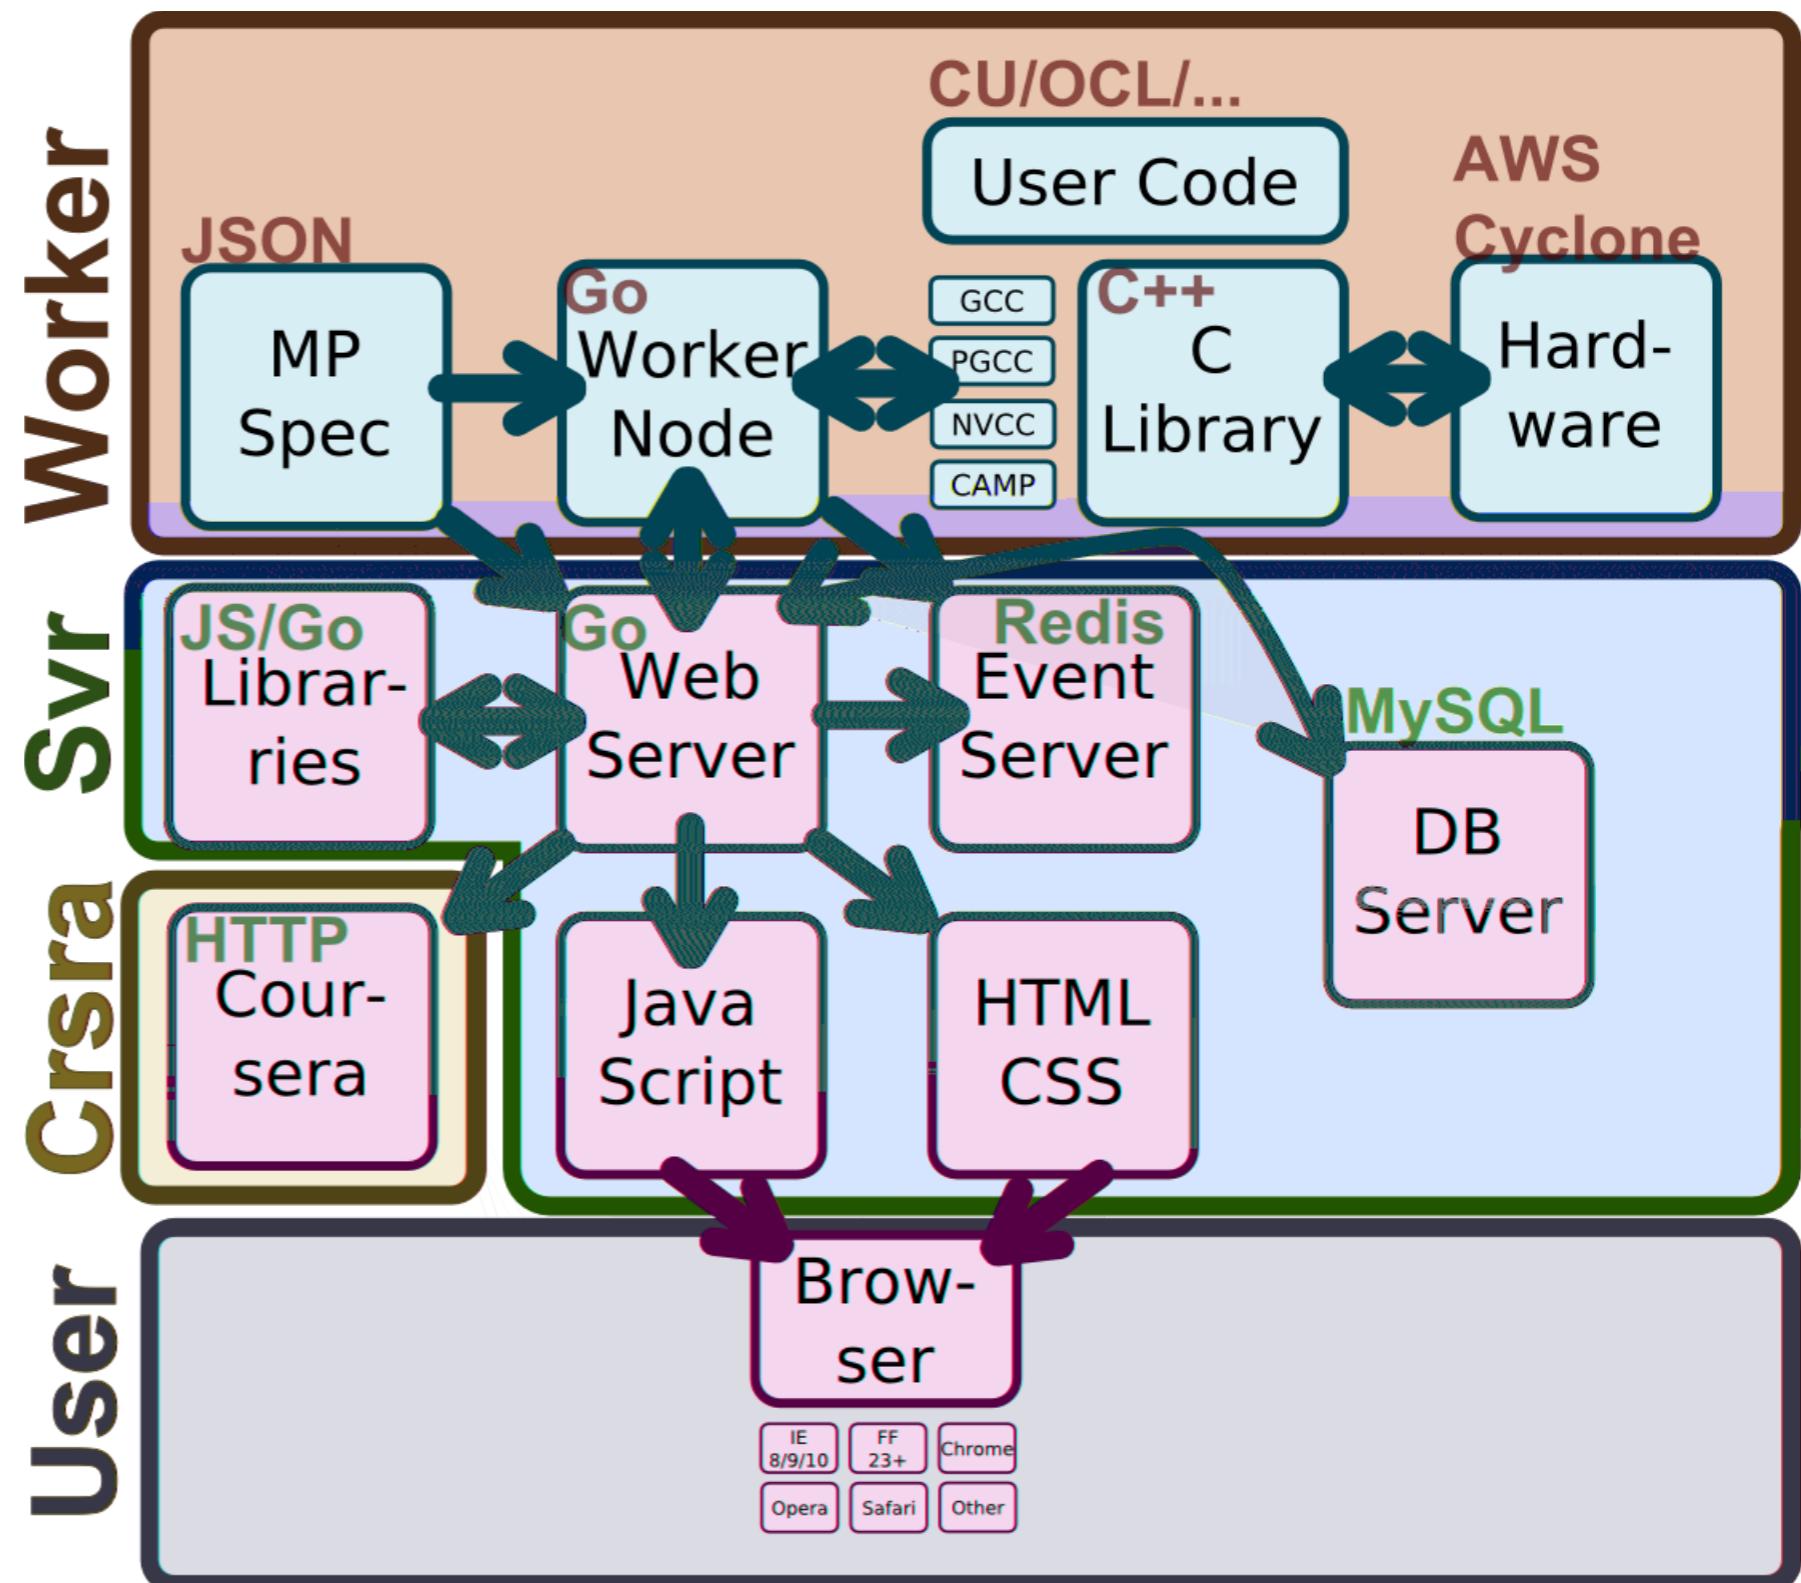
HFRICKOFROSVILLE.COM

Life of a Program Submission

Life of a Program Submission

- Use interface to submit code
- Web-based editor allows changes
- Grade grade center





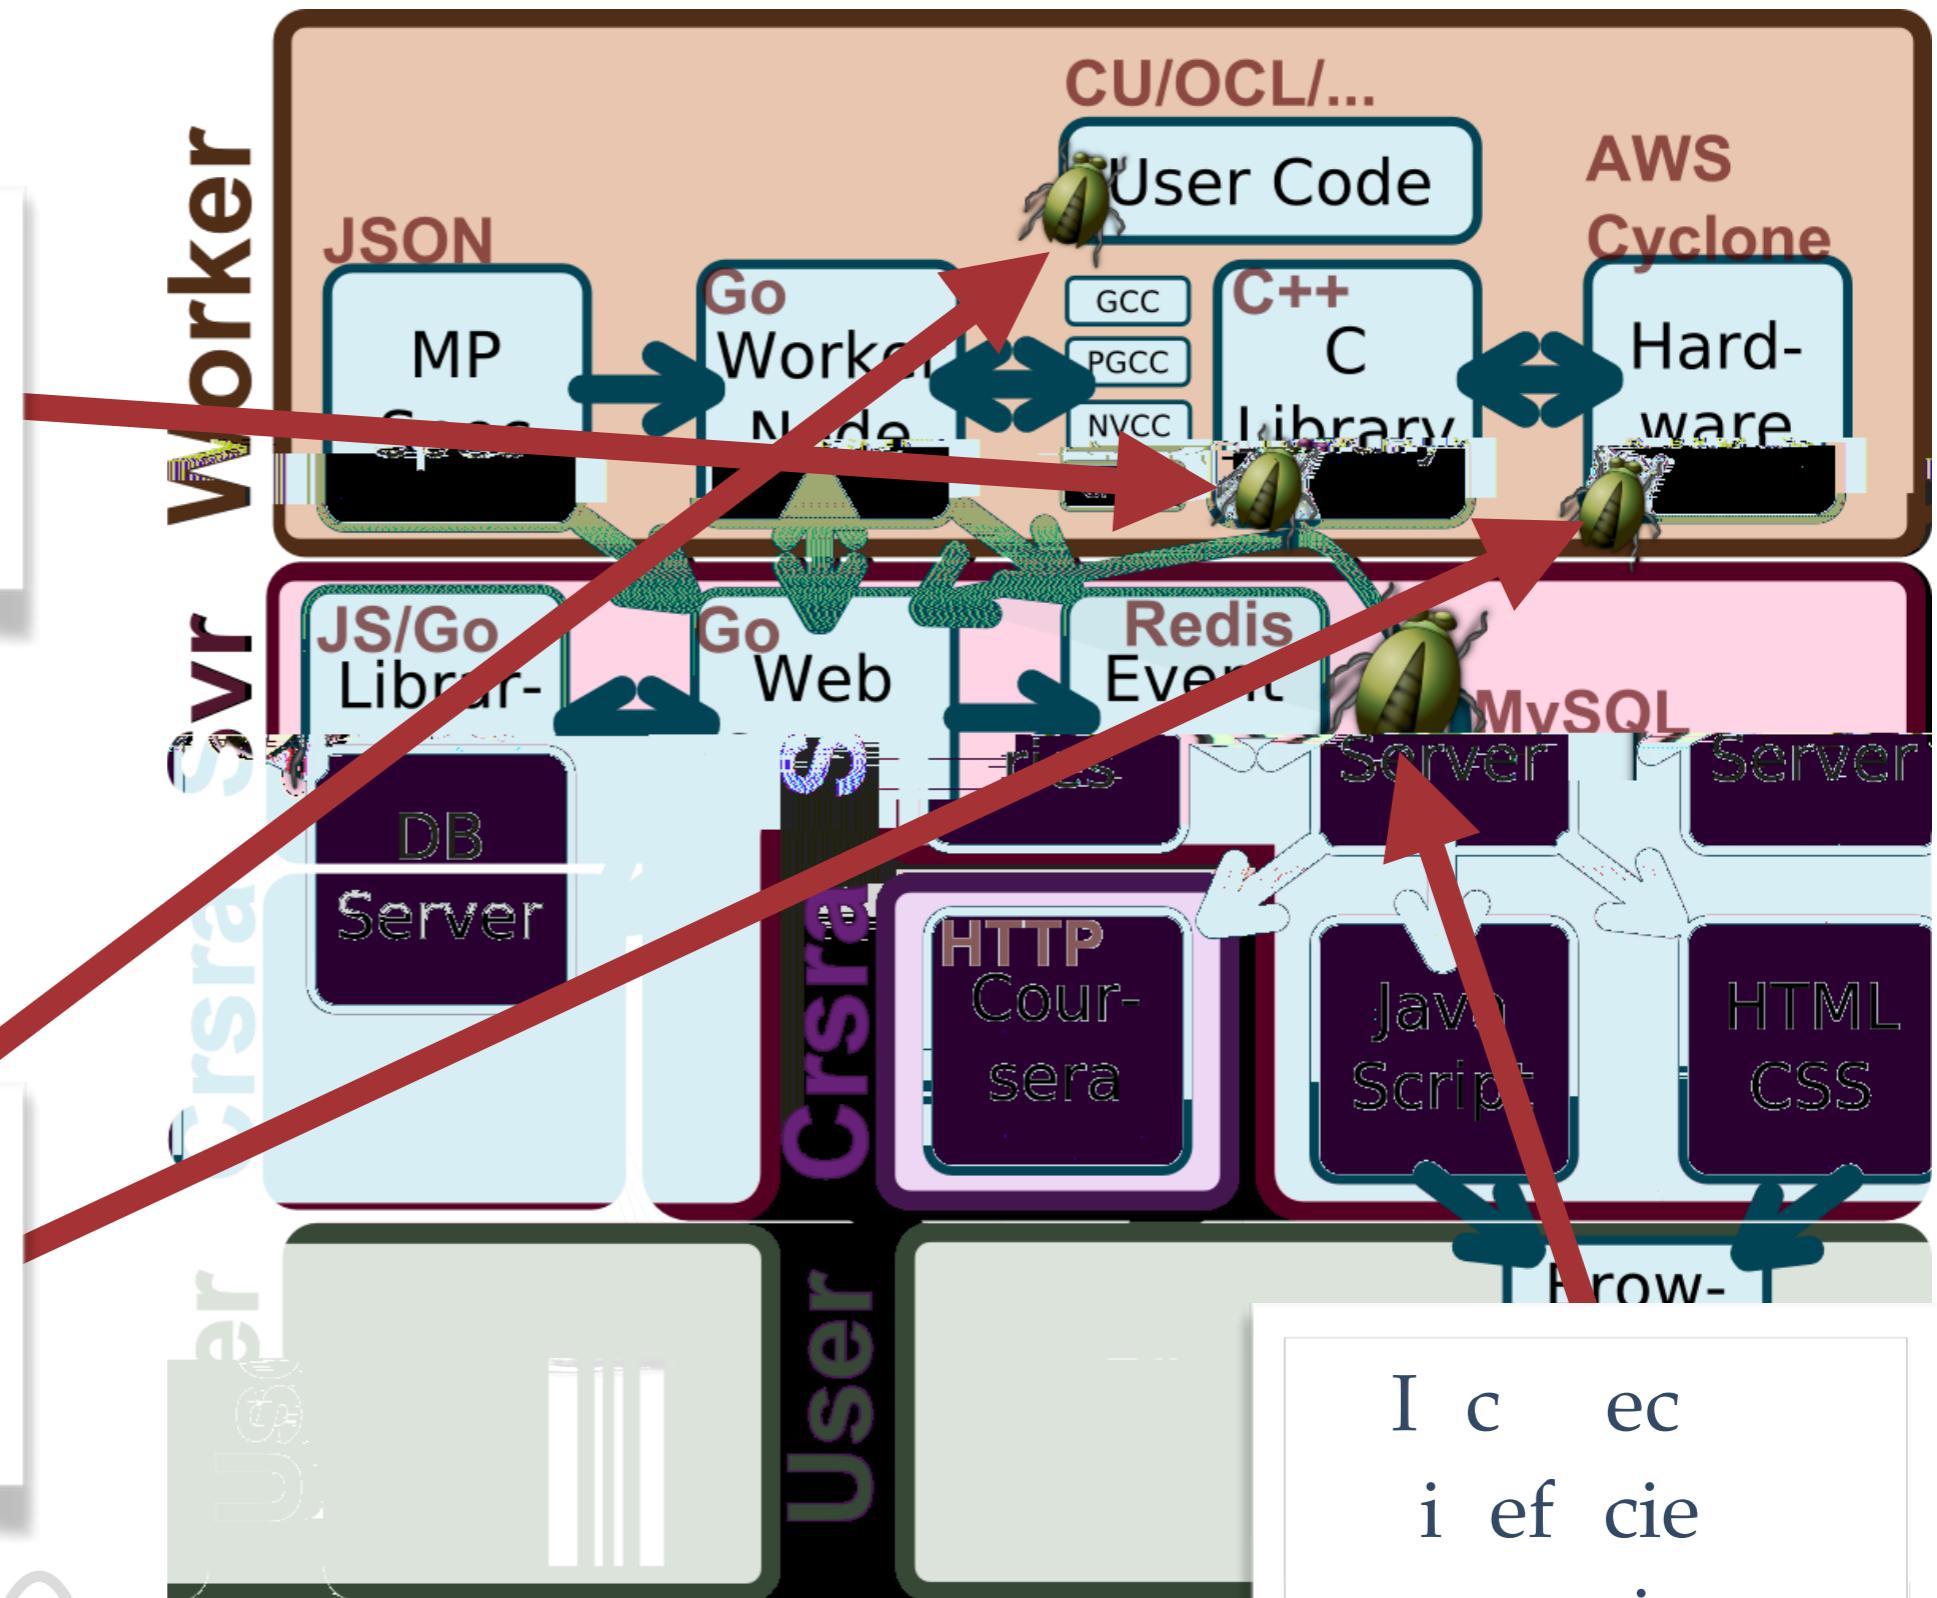
Detailed Architecture

Source of Bugs

Made in
g
f
b ic,
e he
de

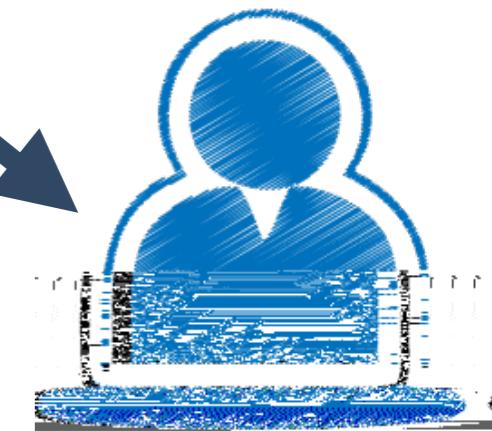
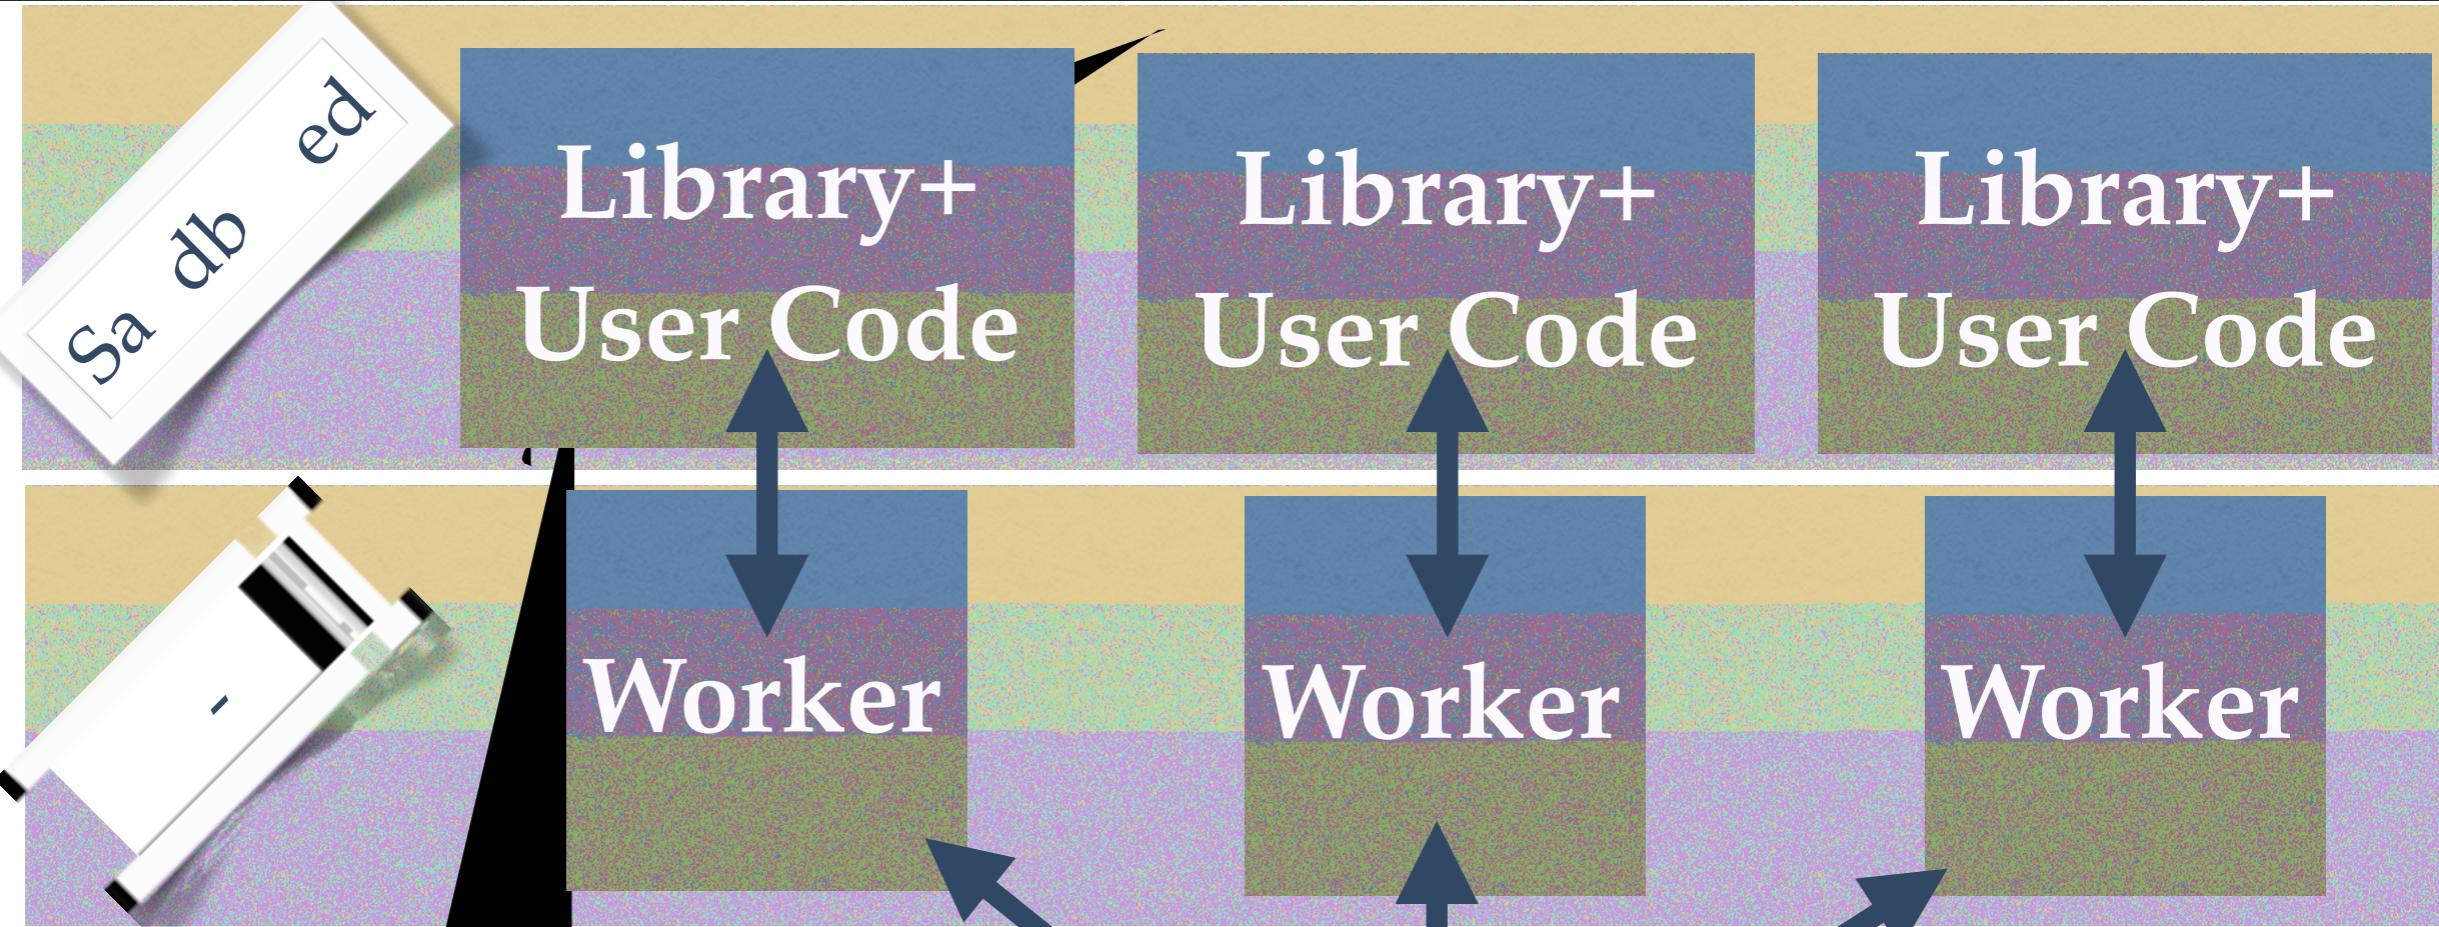
Created a d
ce f a
be (ha e
Cc e a d AWS)

I c ec
i ef cie
e ie



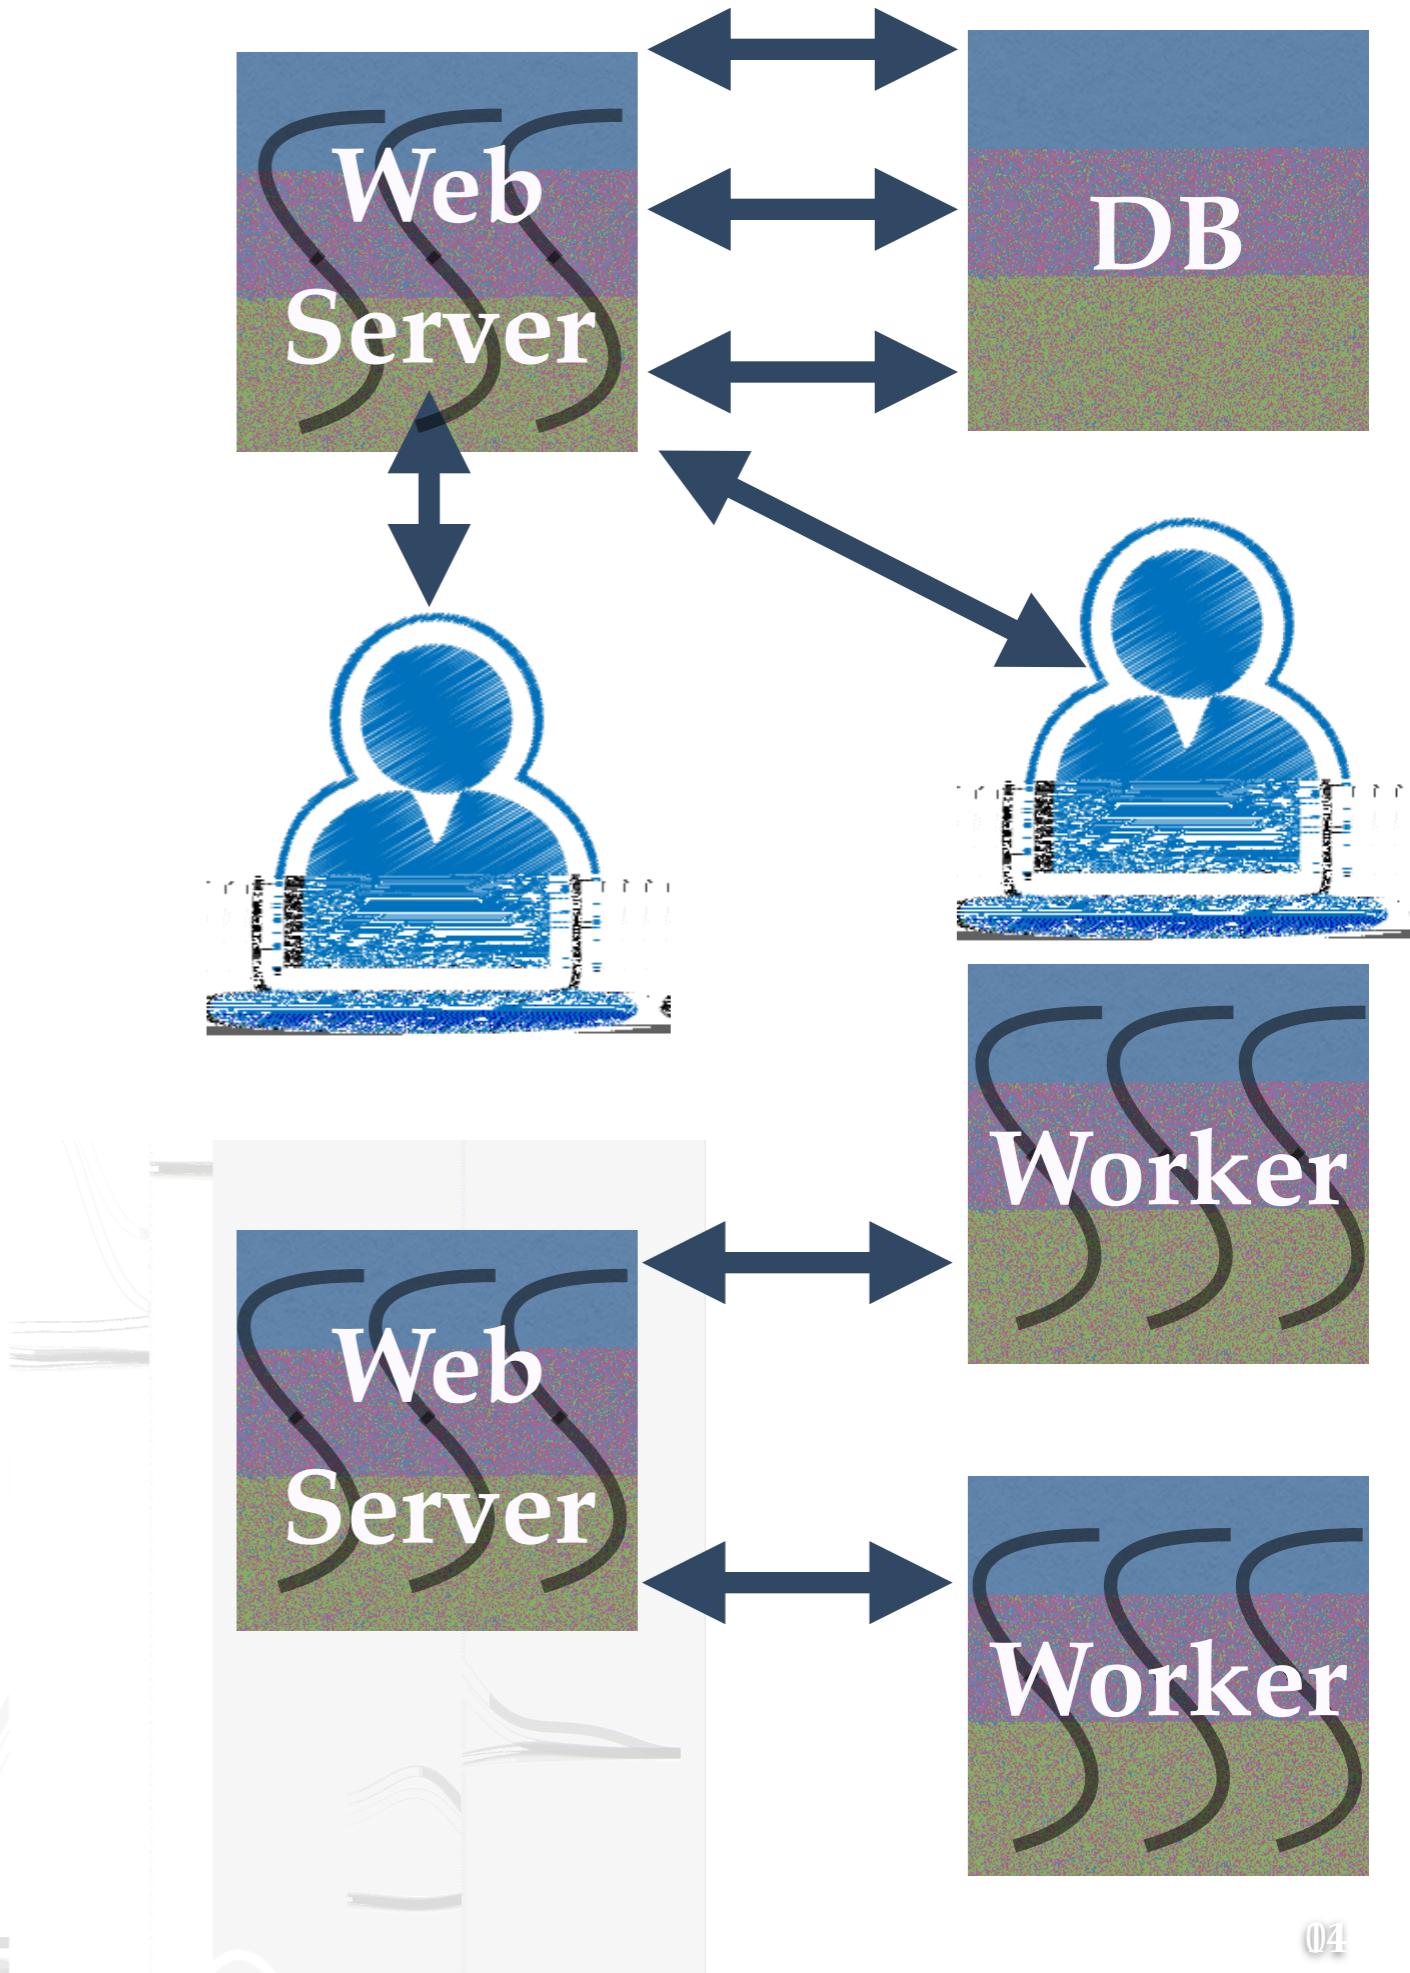
Security

- Sa db e e
- W e , el e
- DB a e
- A i e
80 e e

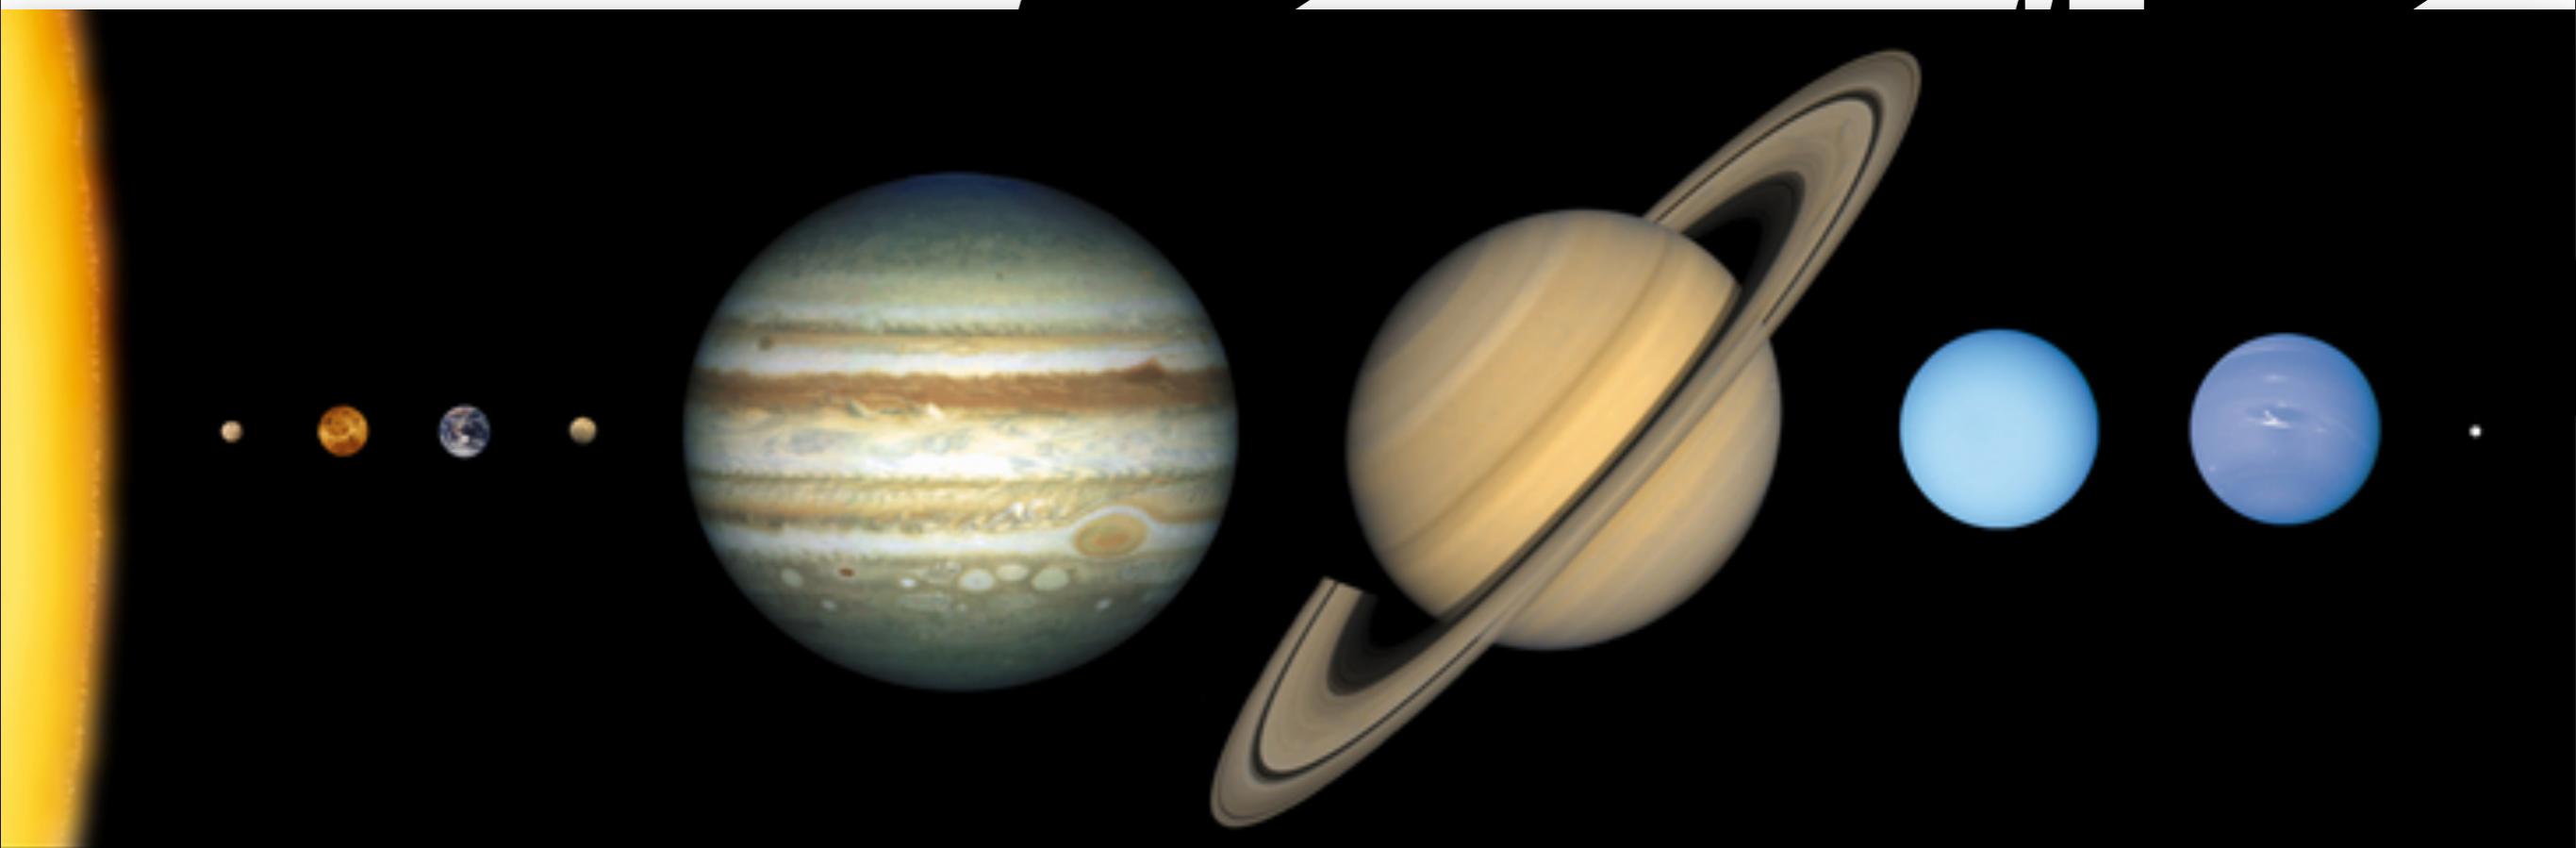


How to Scale?

- Each epoch is a high eight-head (had half the time)
- The high eight heads are added in parallel (being half before each epoch)
- Create each epoch in parallel
- Merge gradients in parallel
- Make each feature parallel
- We can use GPU if available



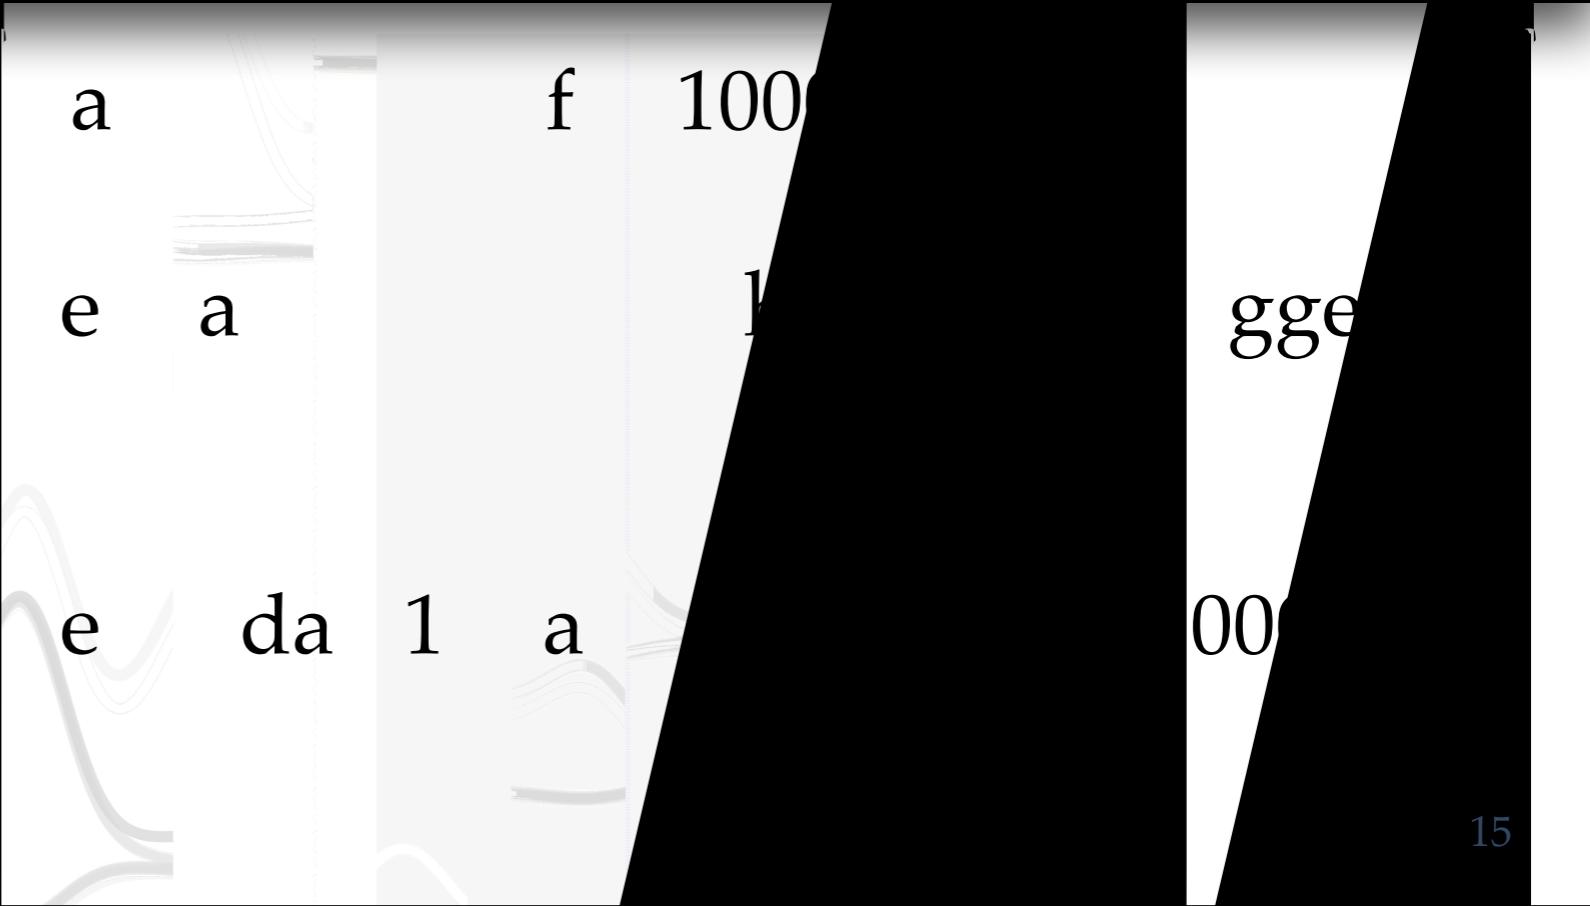
Scale



✿ Wha f 2

✿ Wha f 1
a he a e i e

✿ Wha f
e e da



I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.

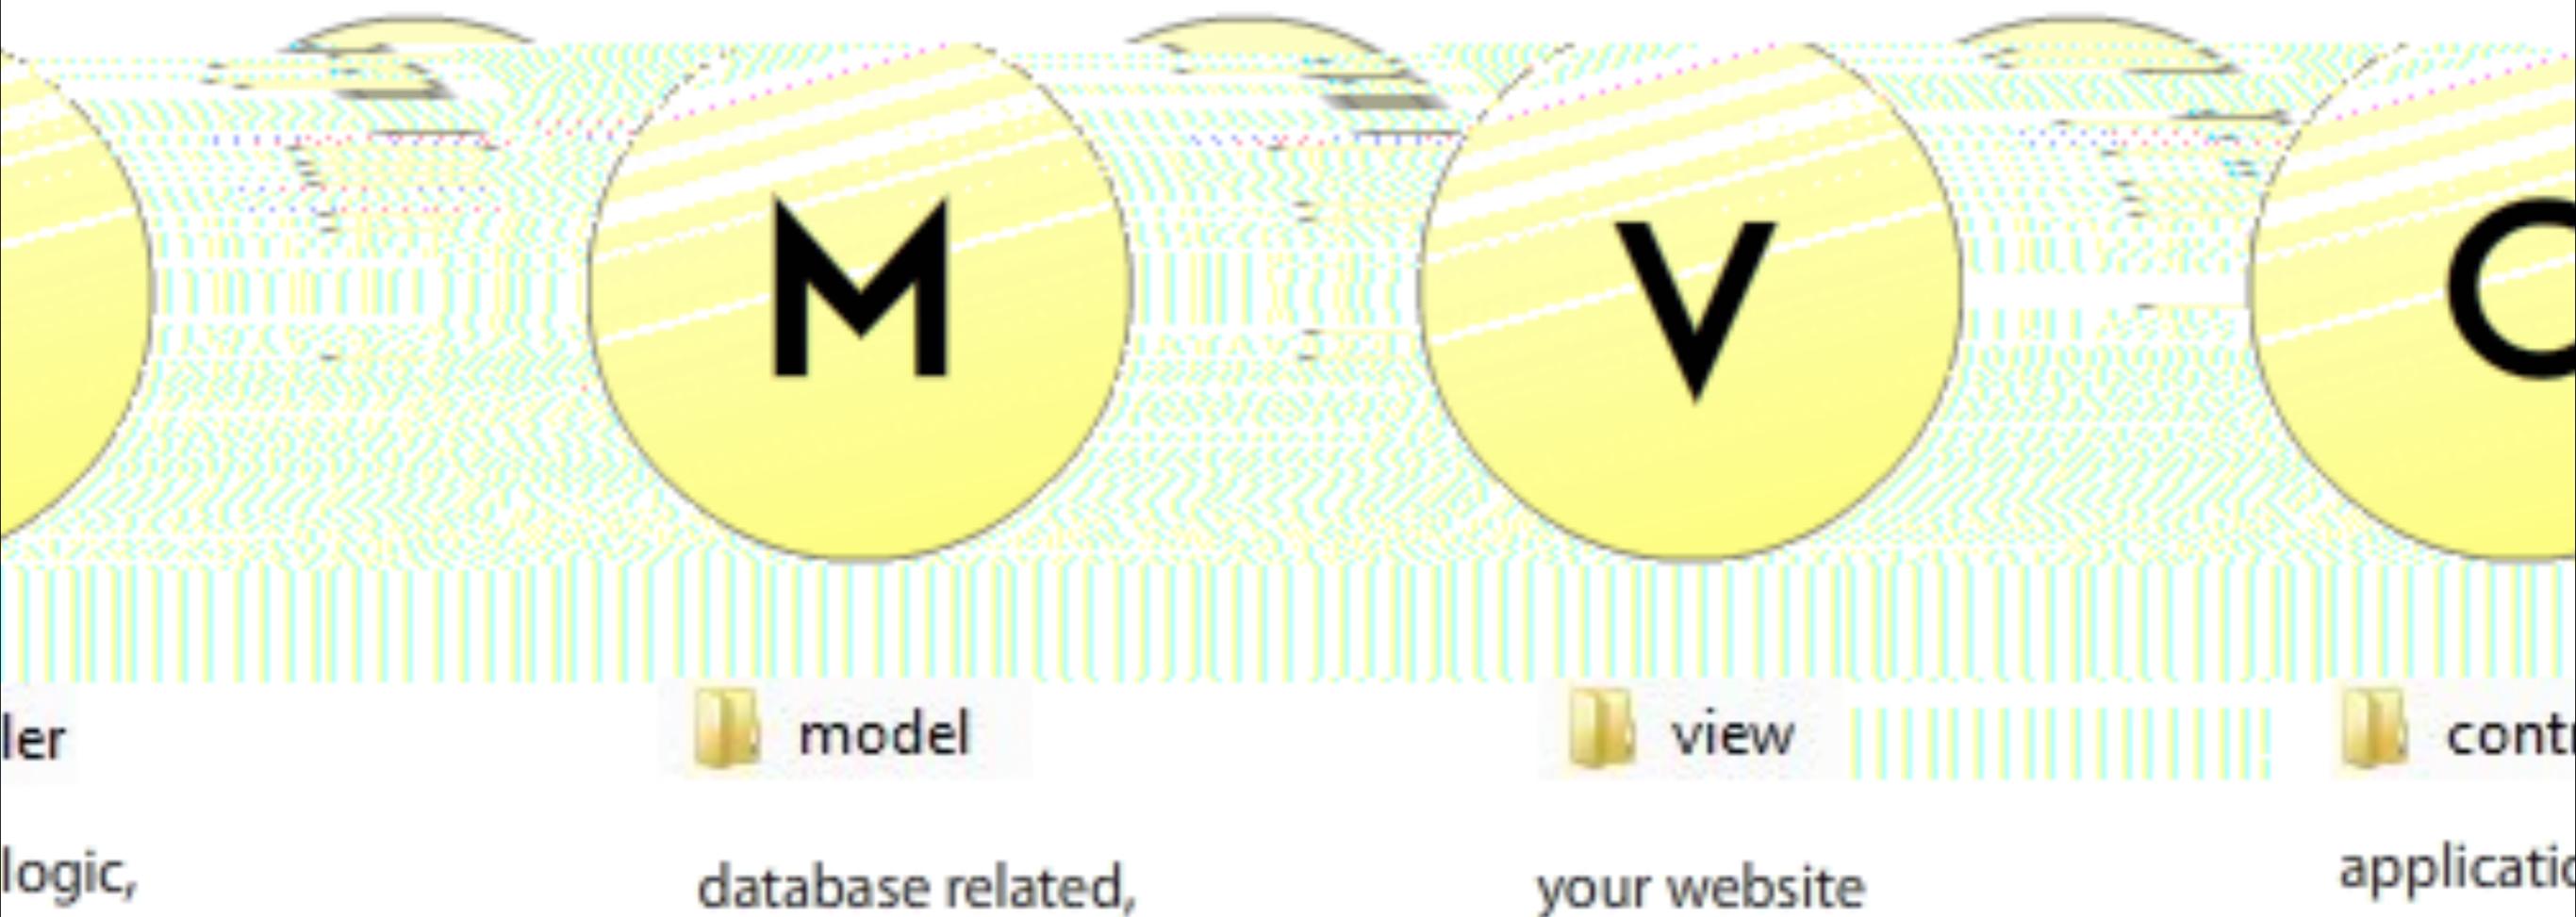


Lessons Learned (Implementation)

Abstraction is not a Good Thing

- * A he a f he c e, e e e i g a ib a
i if da aba e e ie i did i if he , b
ge e a ed c e e ie ha e e ib e
deb g

Abstraction is a Great Thing



- ❖ Whe e e aced he ib a , had e ace he
de c de i he a ica i

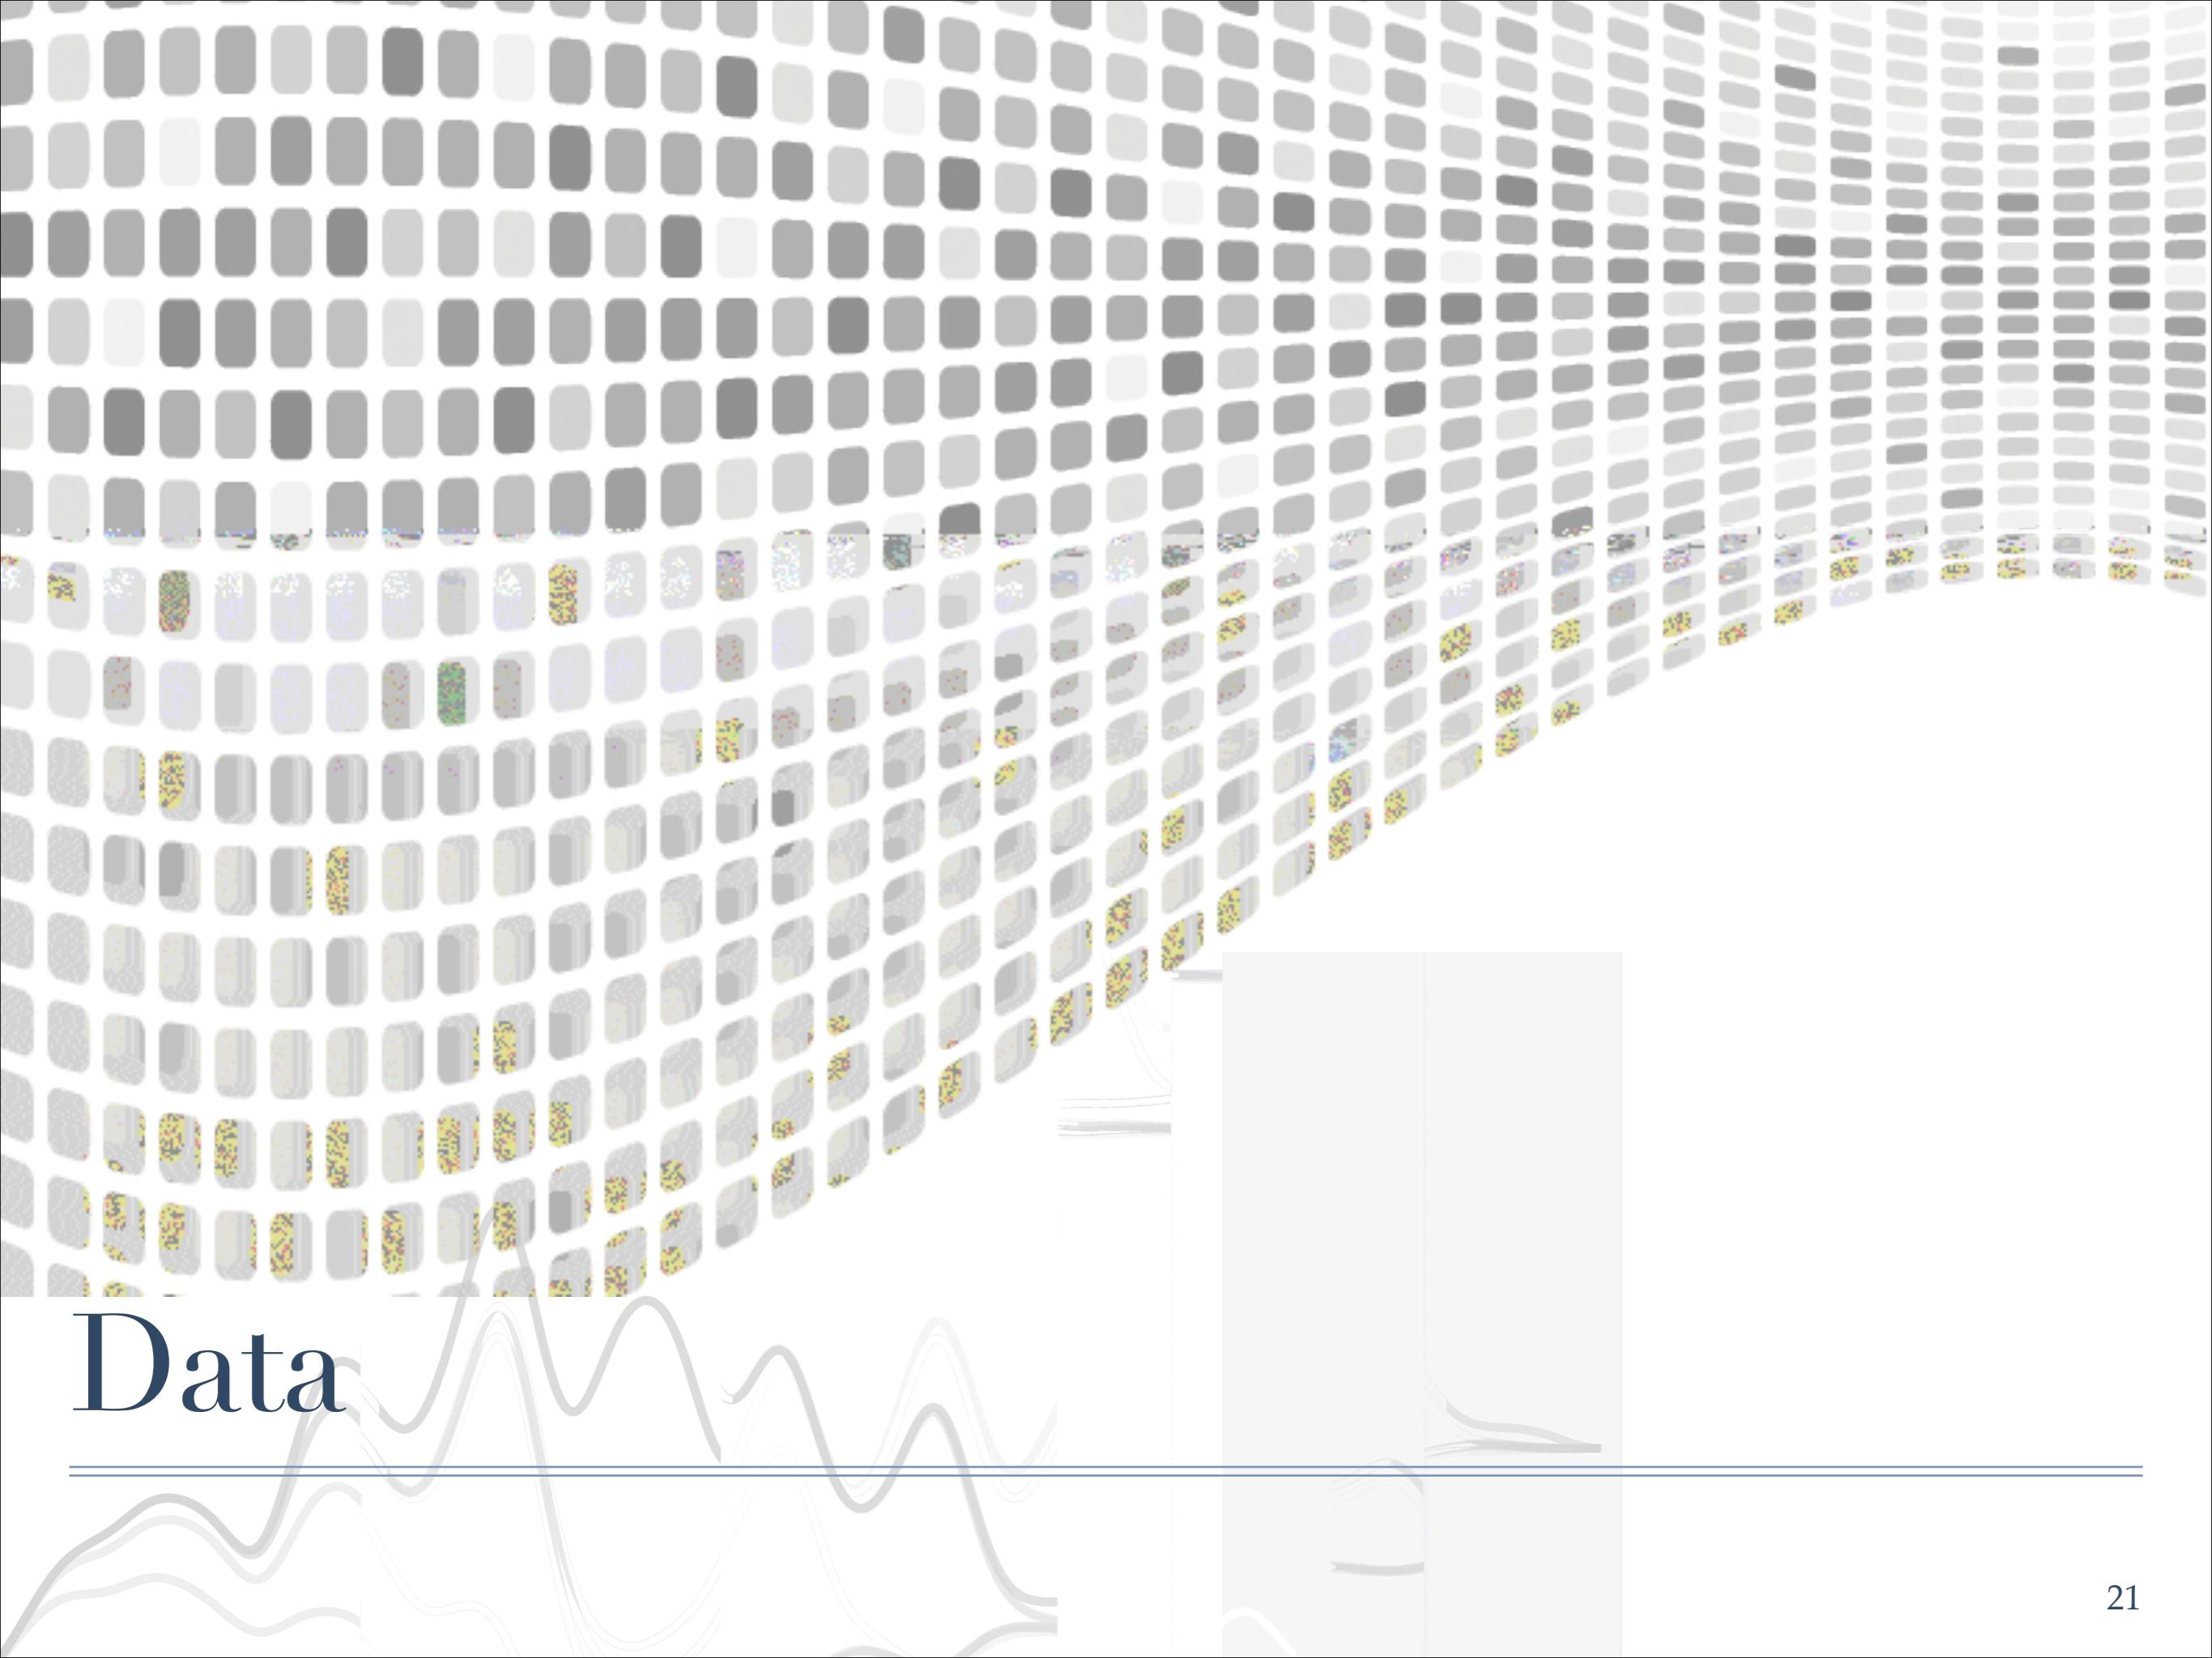
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.



Lessons Learned (Human)

Lessons Learned

- ✿ Negociación es de la vida
- ✿ Y siempre se aprende algo
- ✿ Puedes enseñar a los demás
- ✿ Y siempre se aprende algo más
- ✿ Puedes inspirar a los demás
- ✿ The siempre se aprende algo



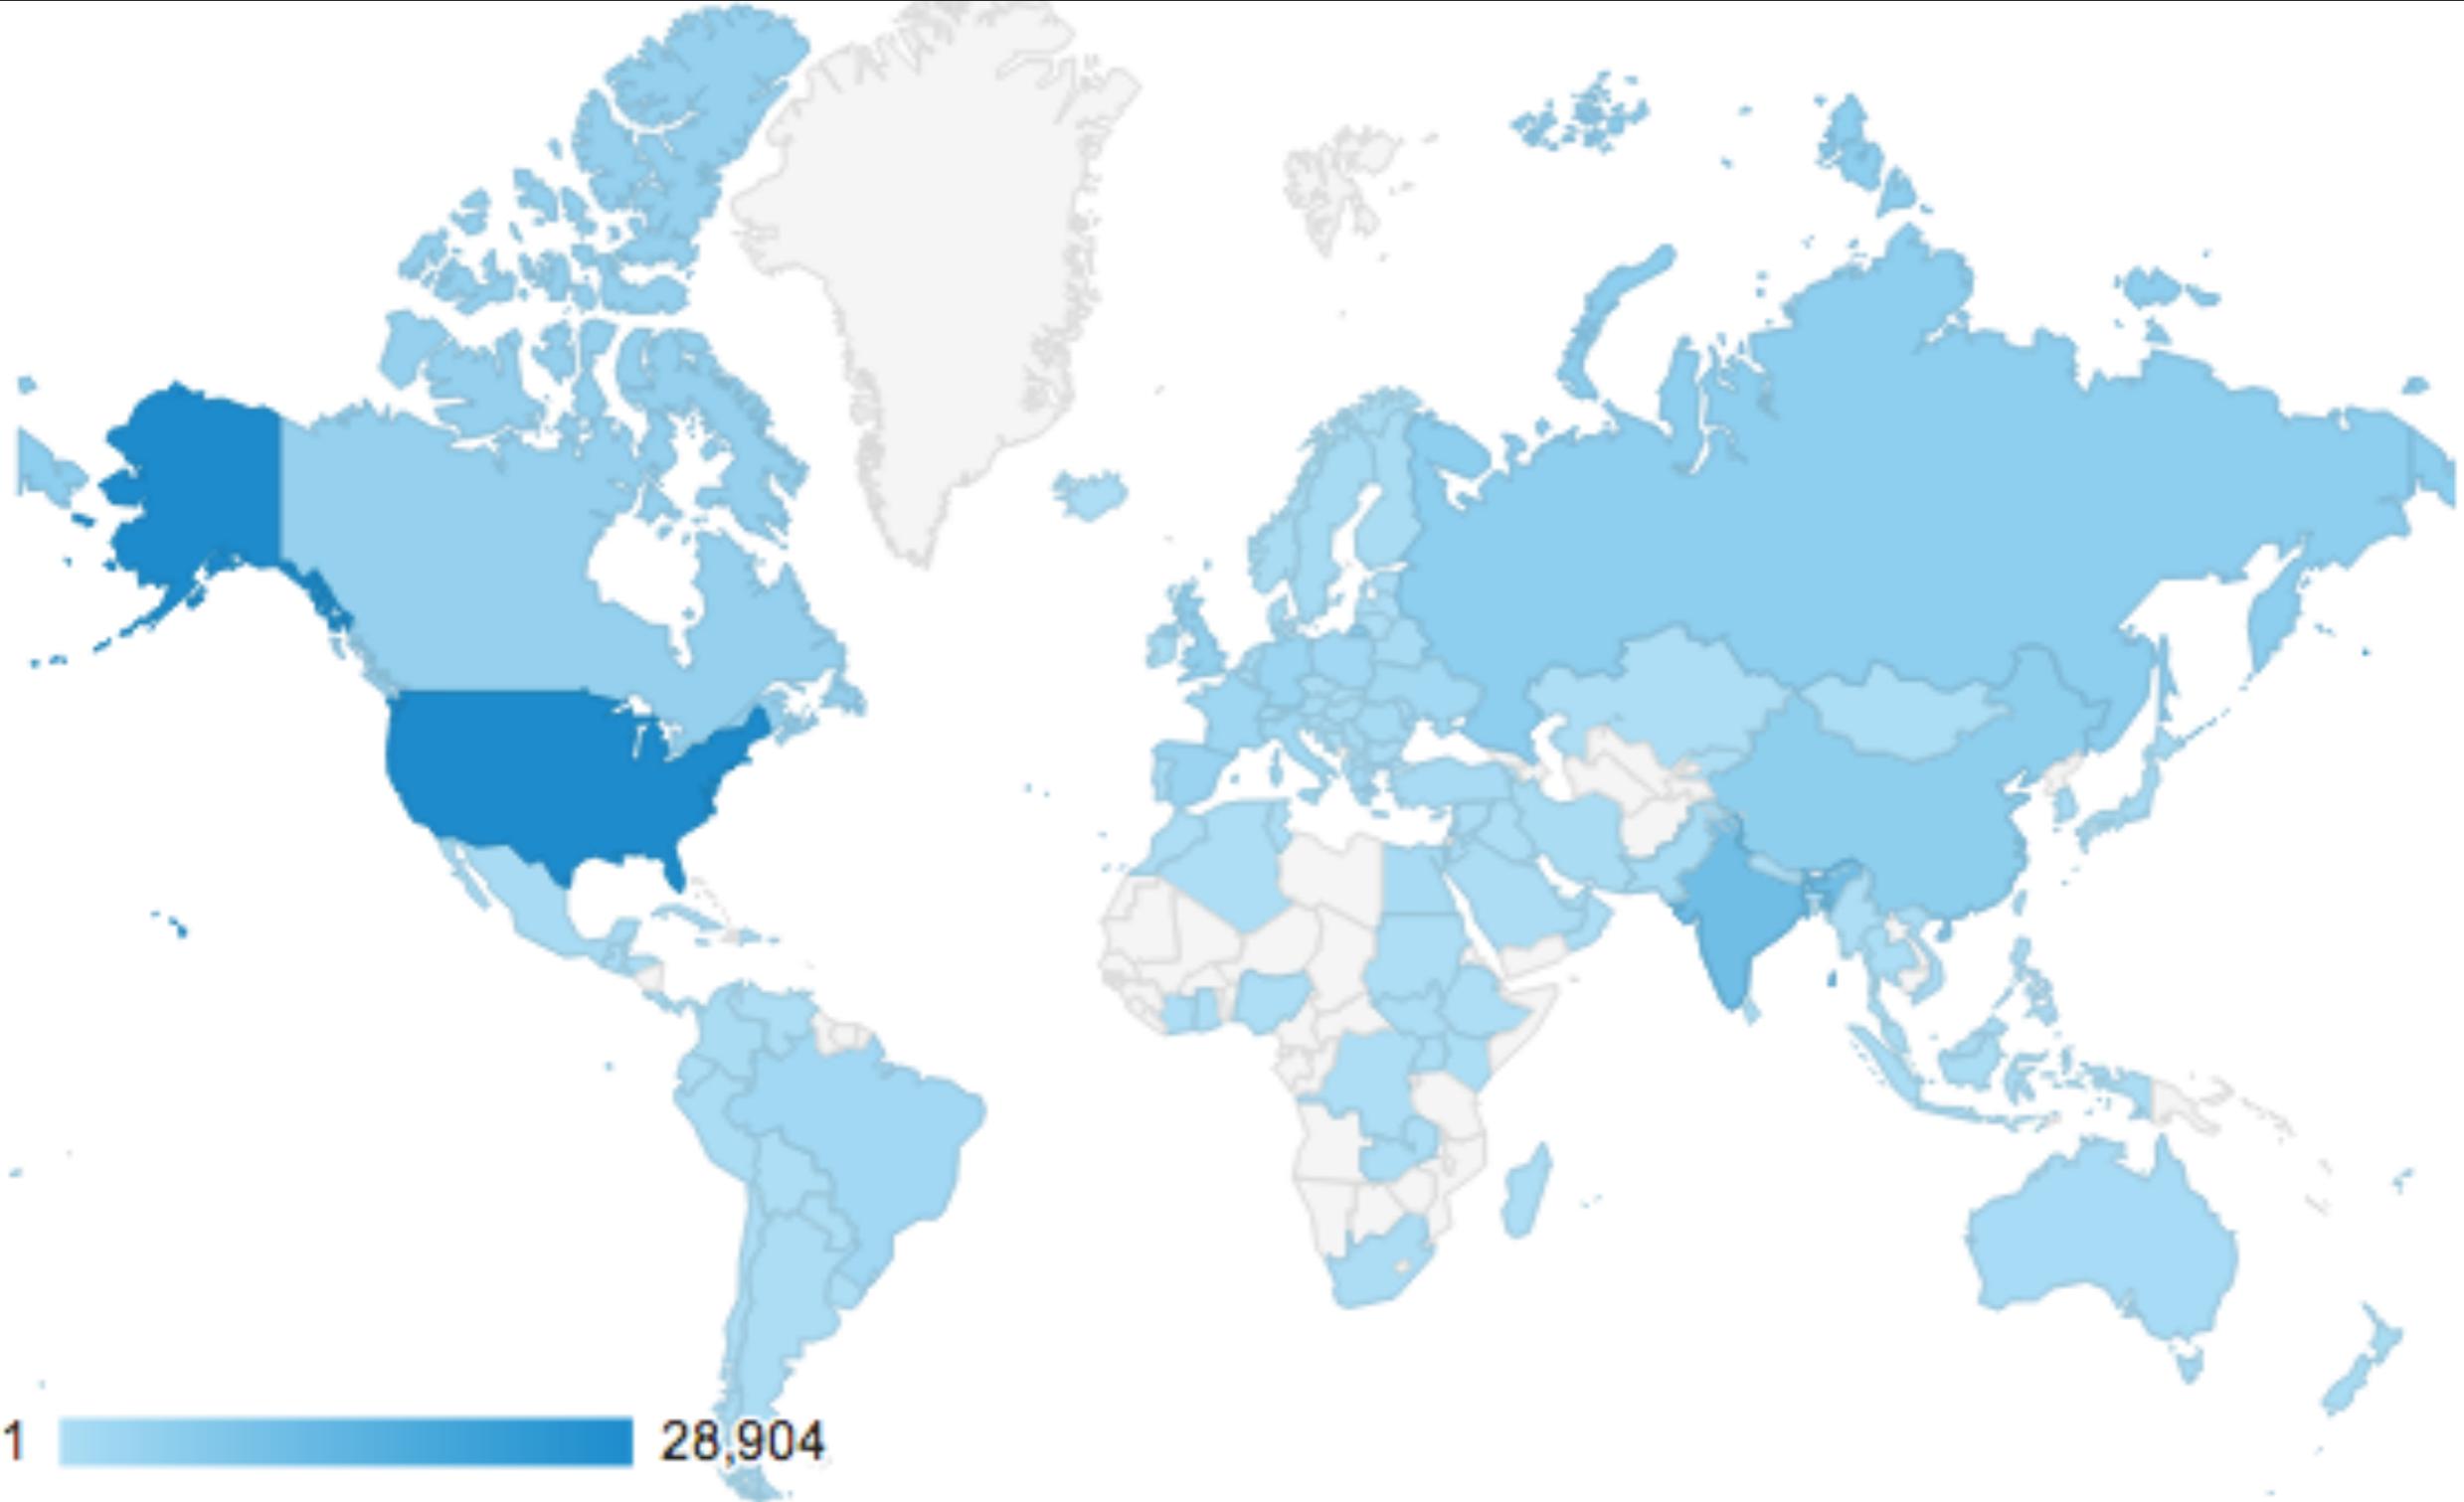
Data

Data Collected

- G geA a ic
- A 13Gb da aba ec ai i g
- 2 Mi i g a e i i
- 700 h a d g a
- R i e/C ia i i e / e
- 6 h a d g aded g a
- A 7Gb E e da aba ec ai i g
- CPU/GPU if ai
- Wha age e Vi i ed



Thousands of Visitors



Users from All over the World

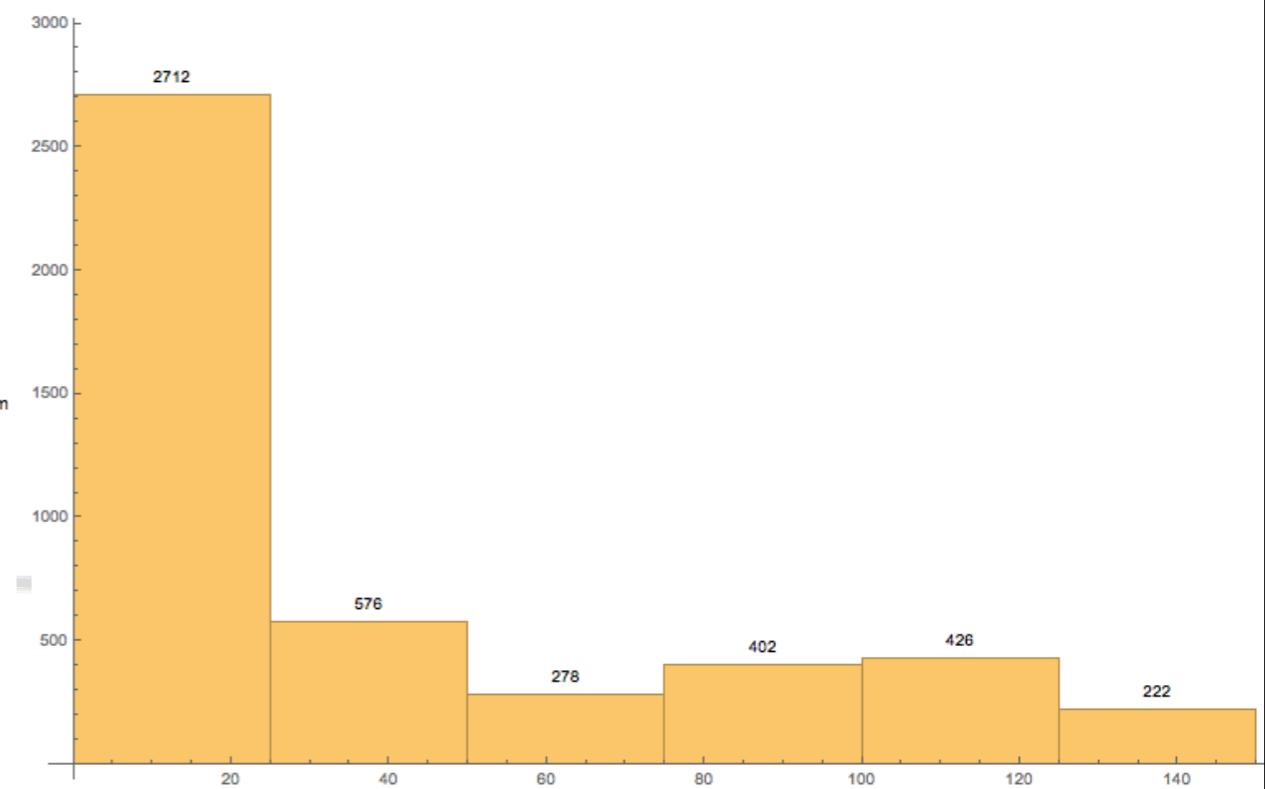
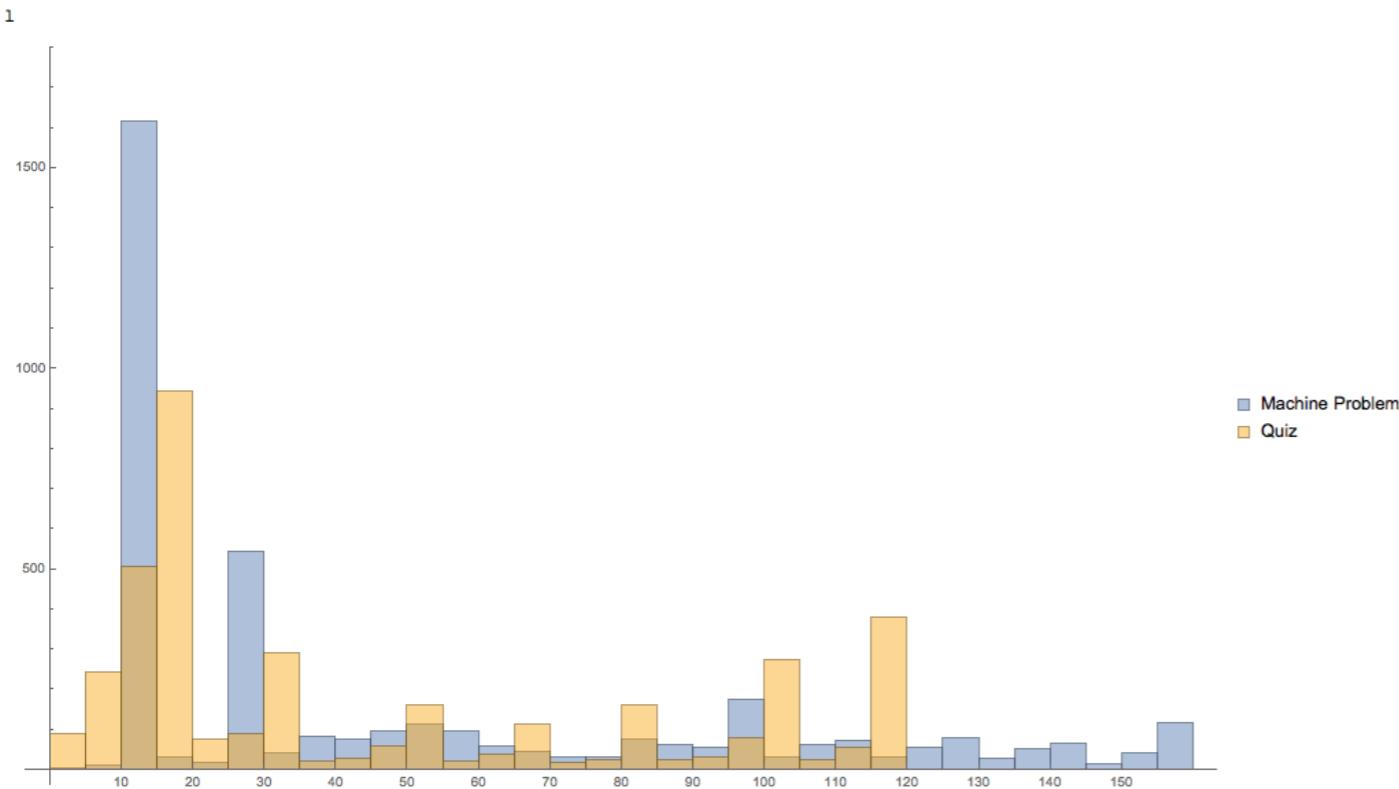


People Spend Time on the Site



Grades

Most People Passed



✿ 939 ac a a ed he c e XXXX hi i c ec FIXME

✿ 989 g e 80%

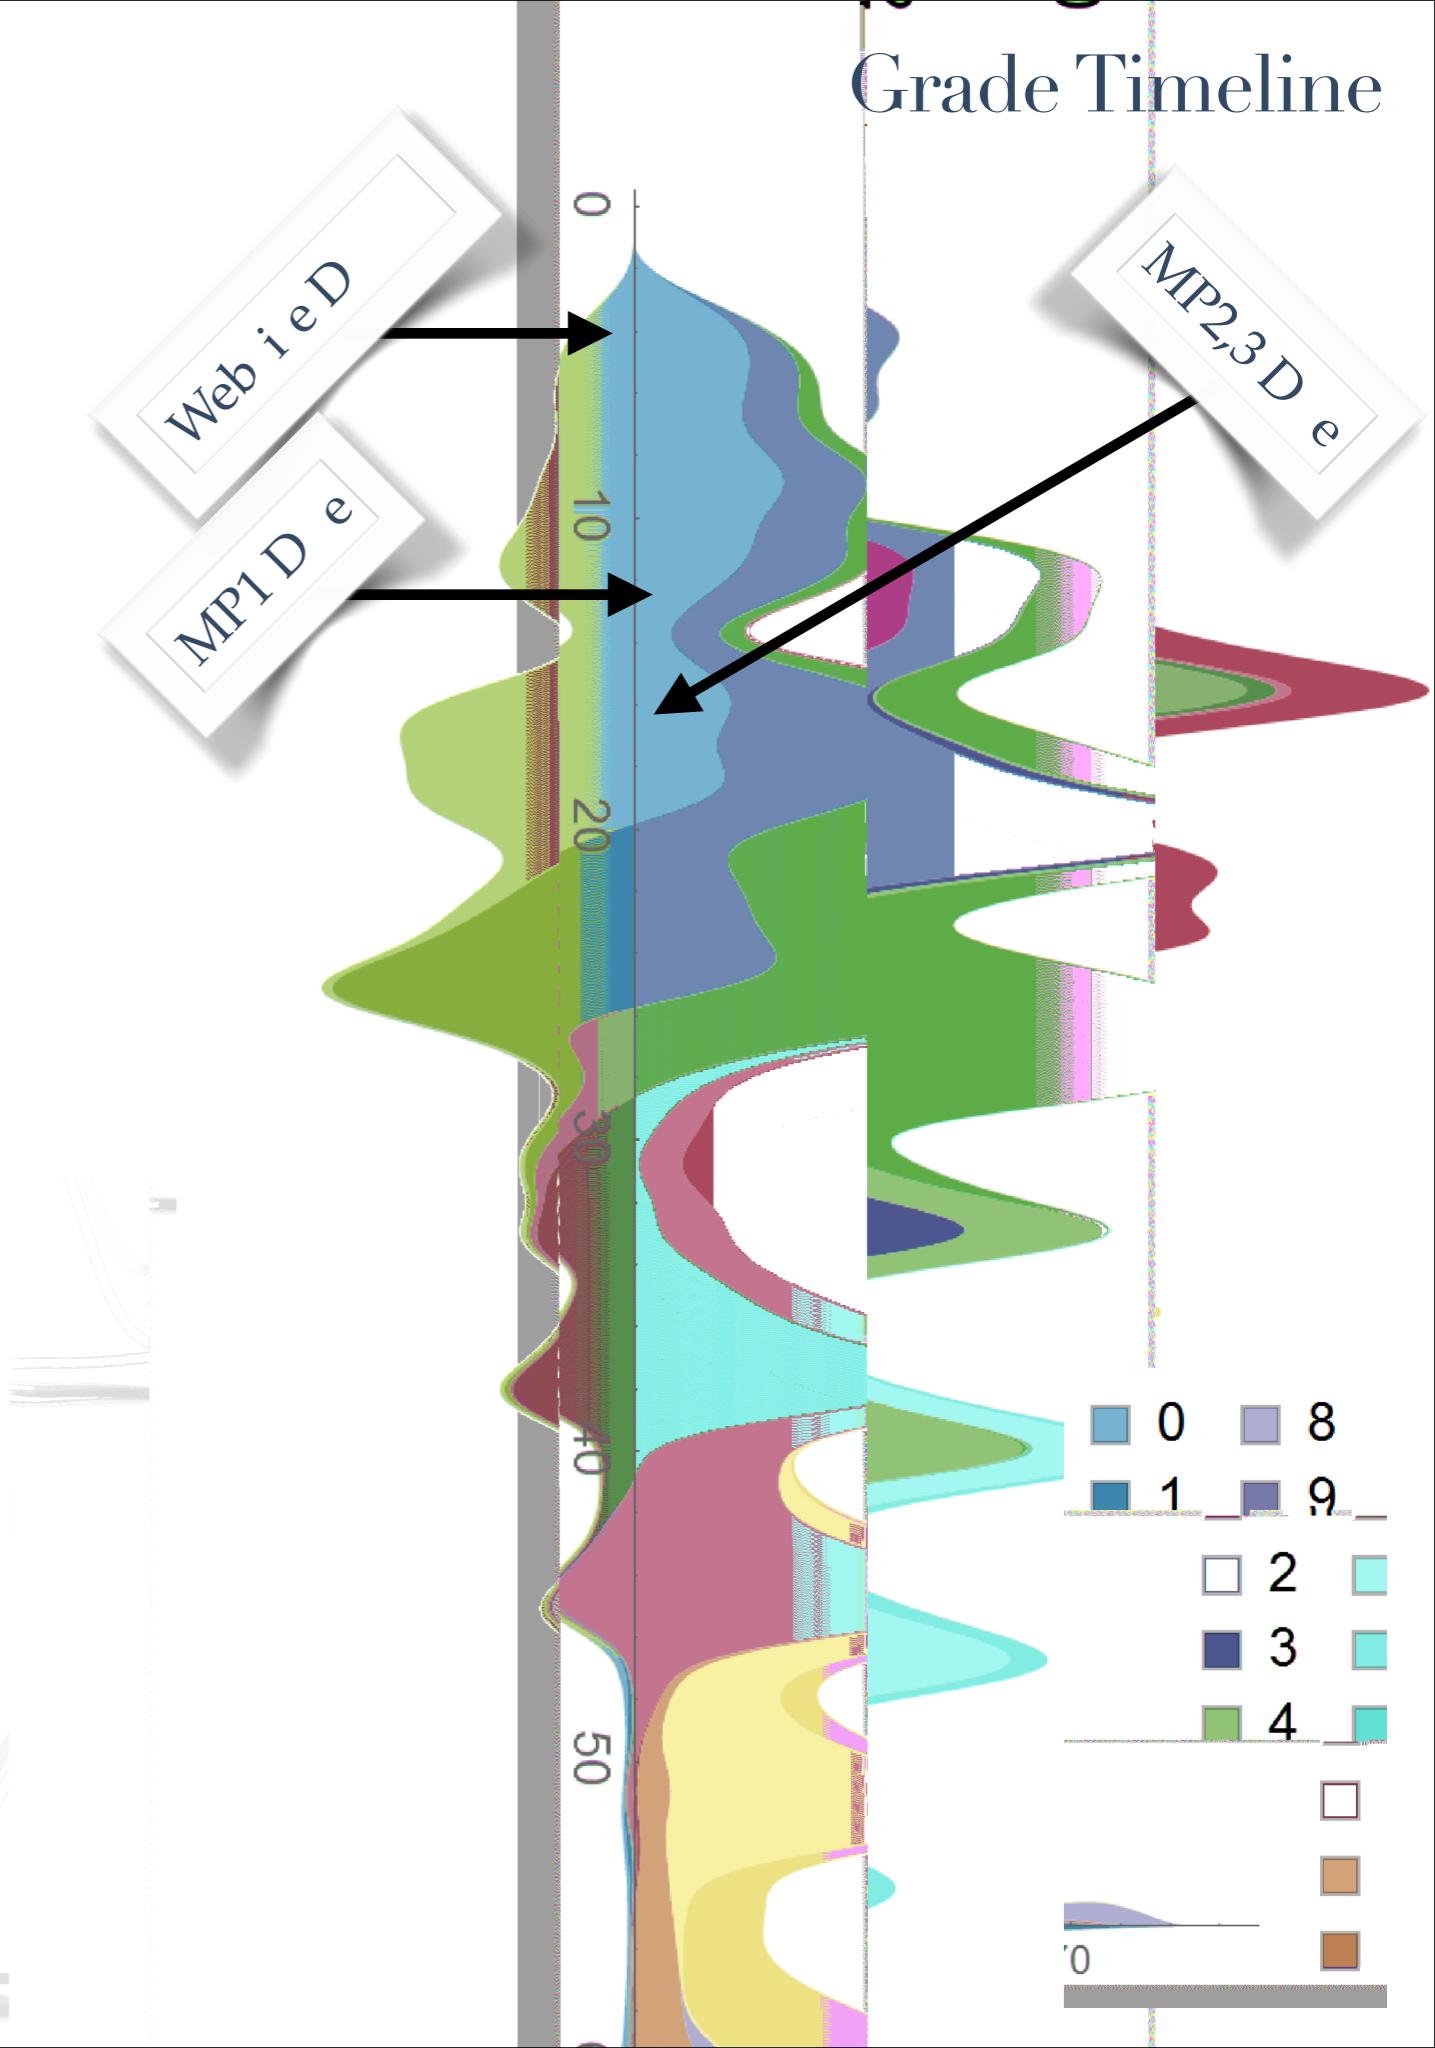
✿ 648 g e 100%

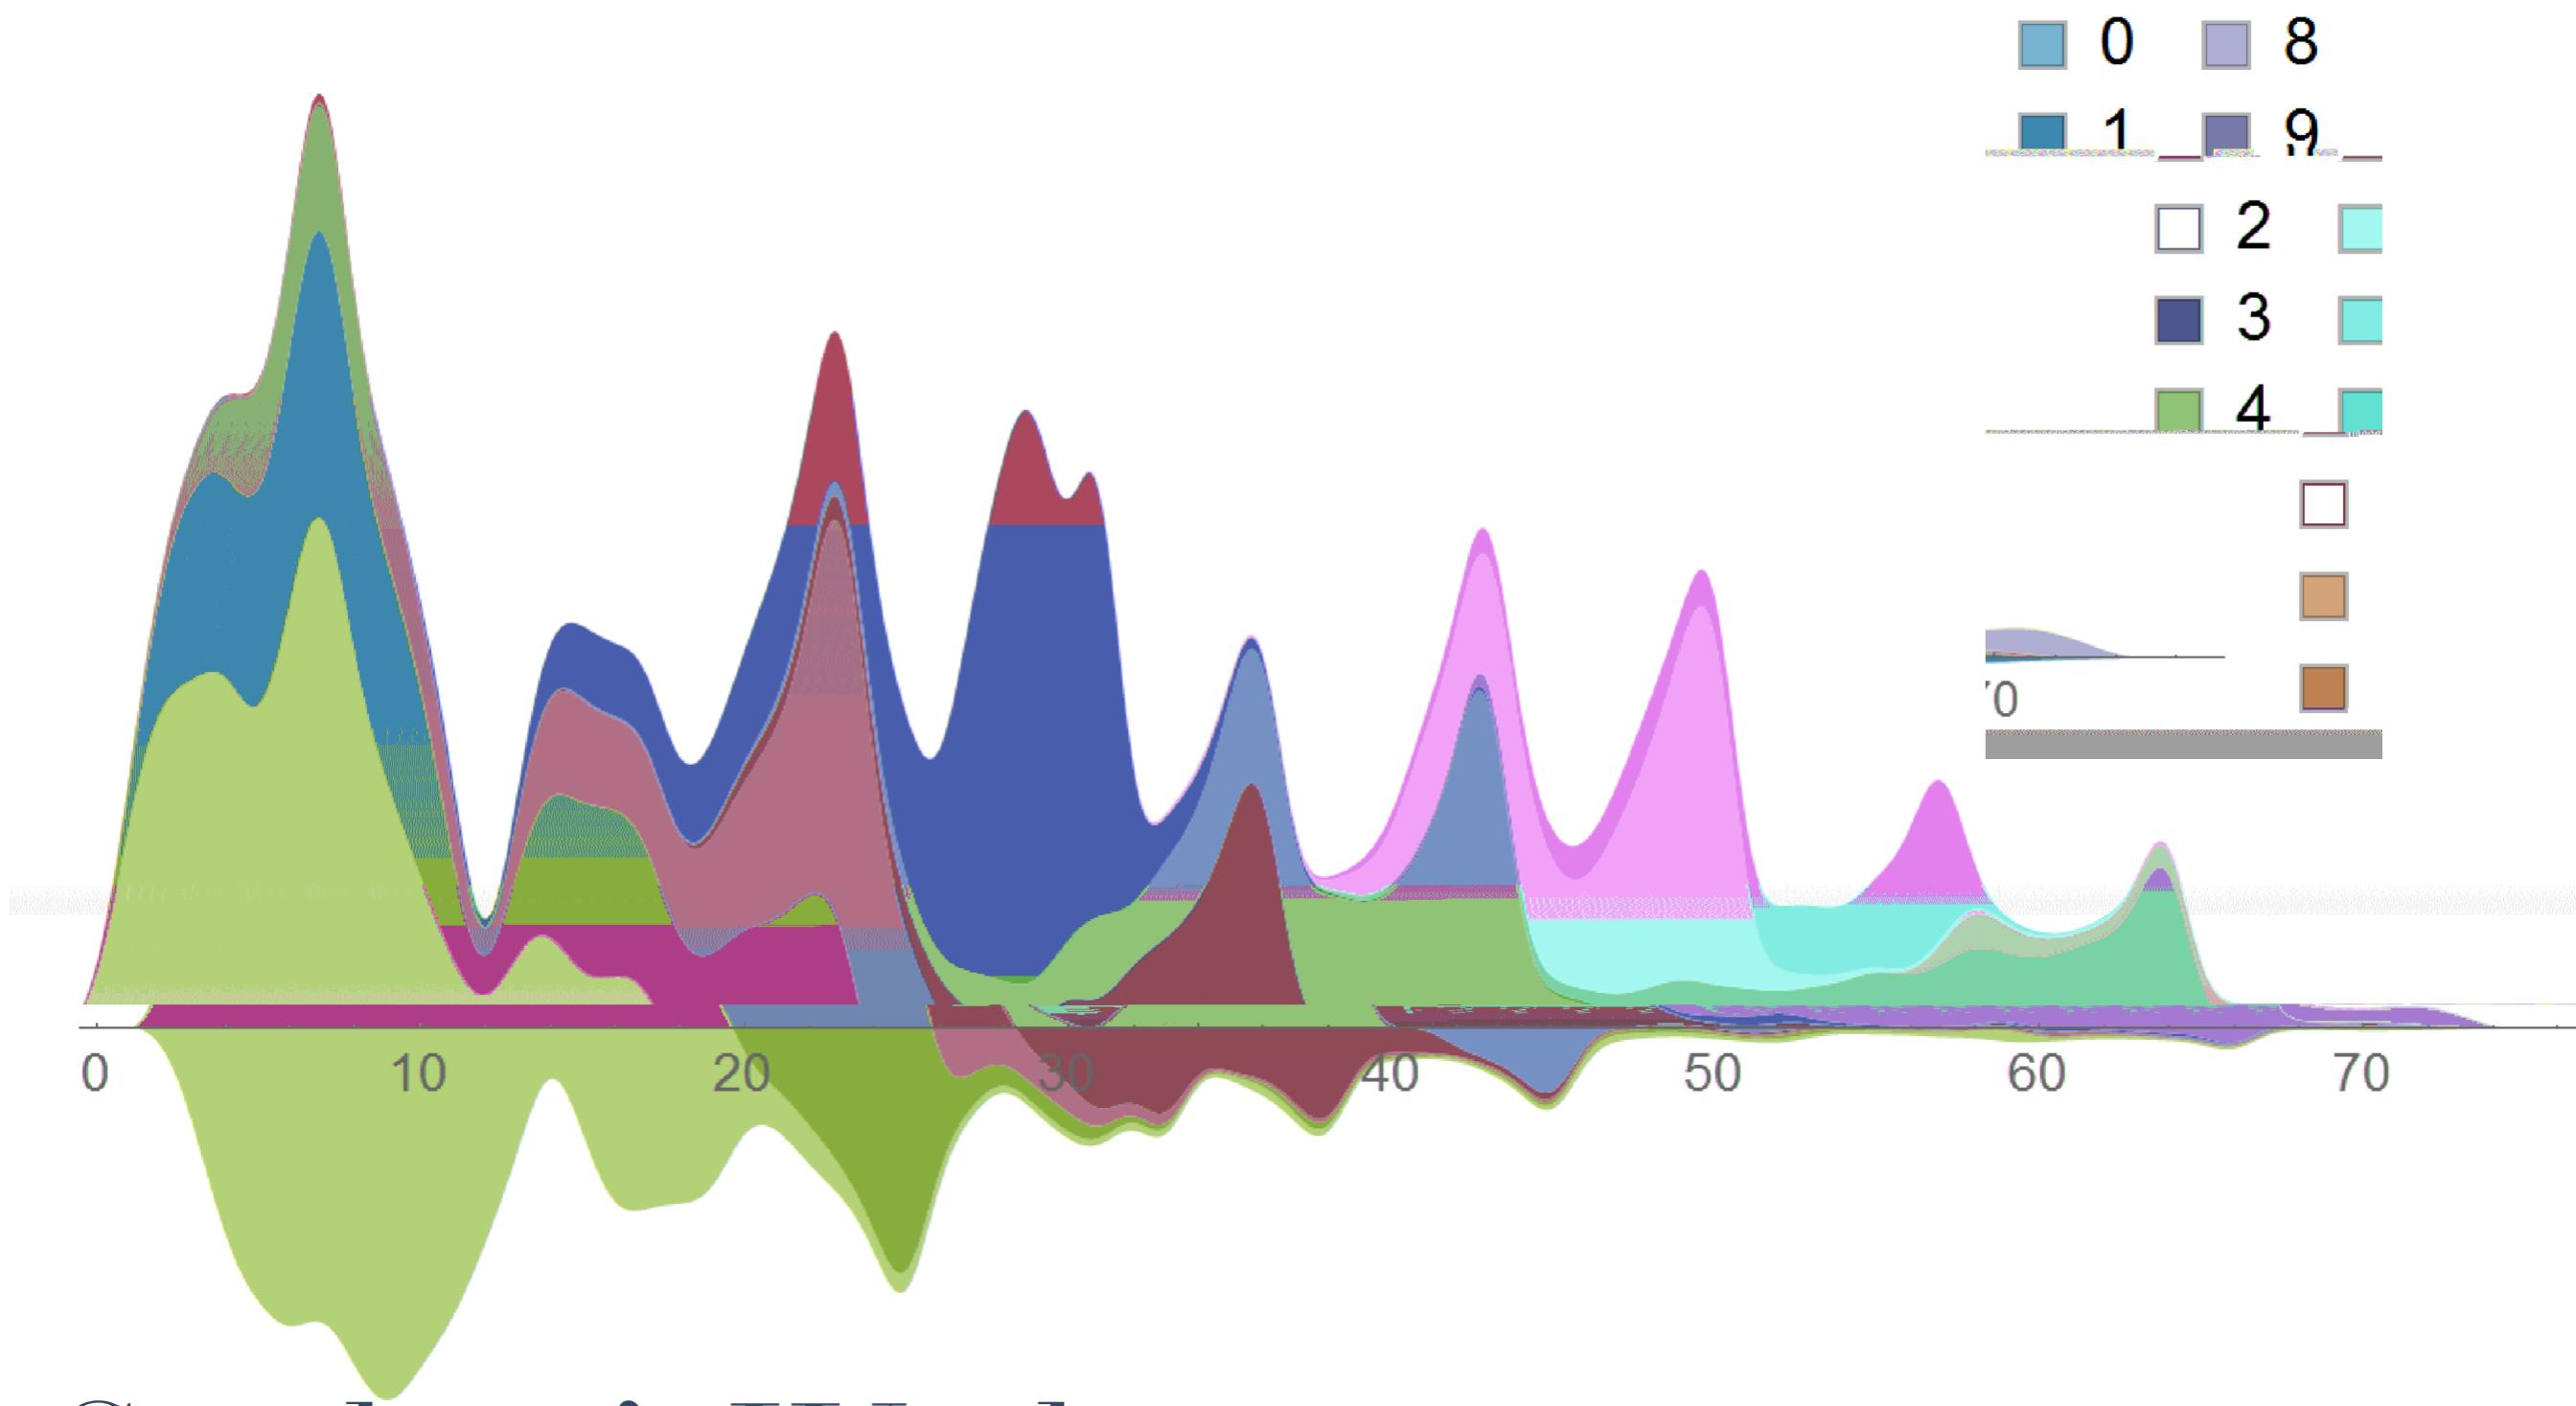
Some Data Insights

Date

People Submit at the Last Minute

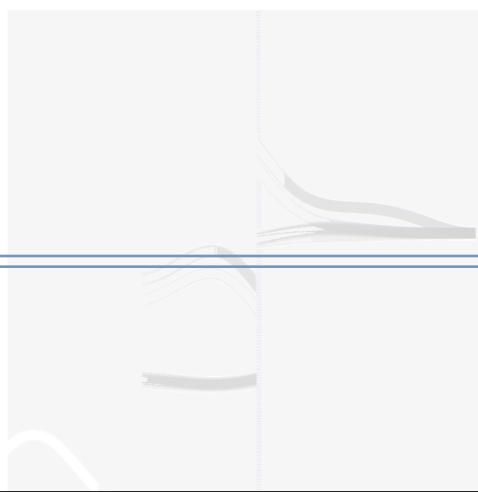
- We ade a c e MP
a ai ab e da 1
- E e e ai ed i he a
i e
- E e MP ha 2 b e
f b i i g he c de a d
a he f he ee e ie





Student's Work

P g a a e i e i e

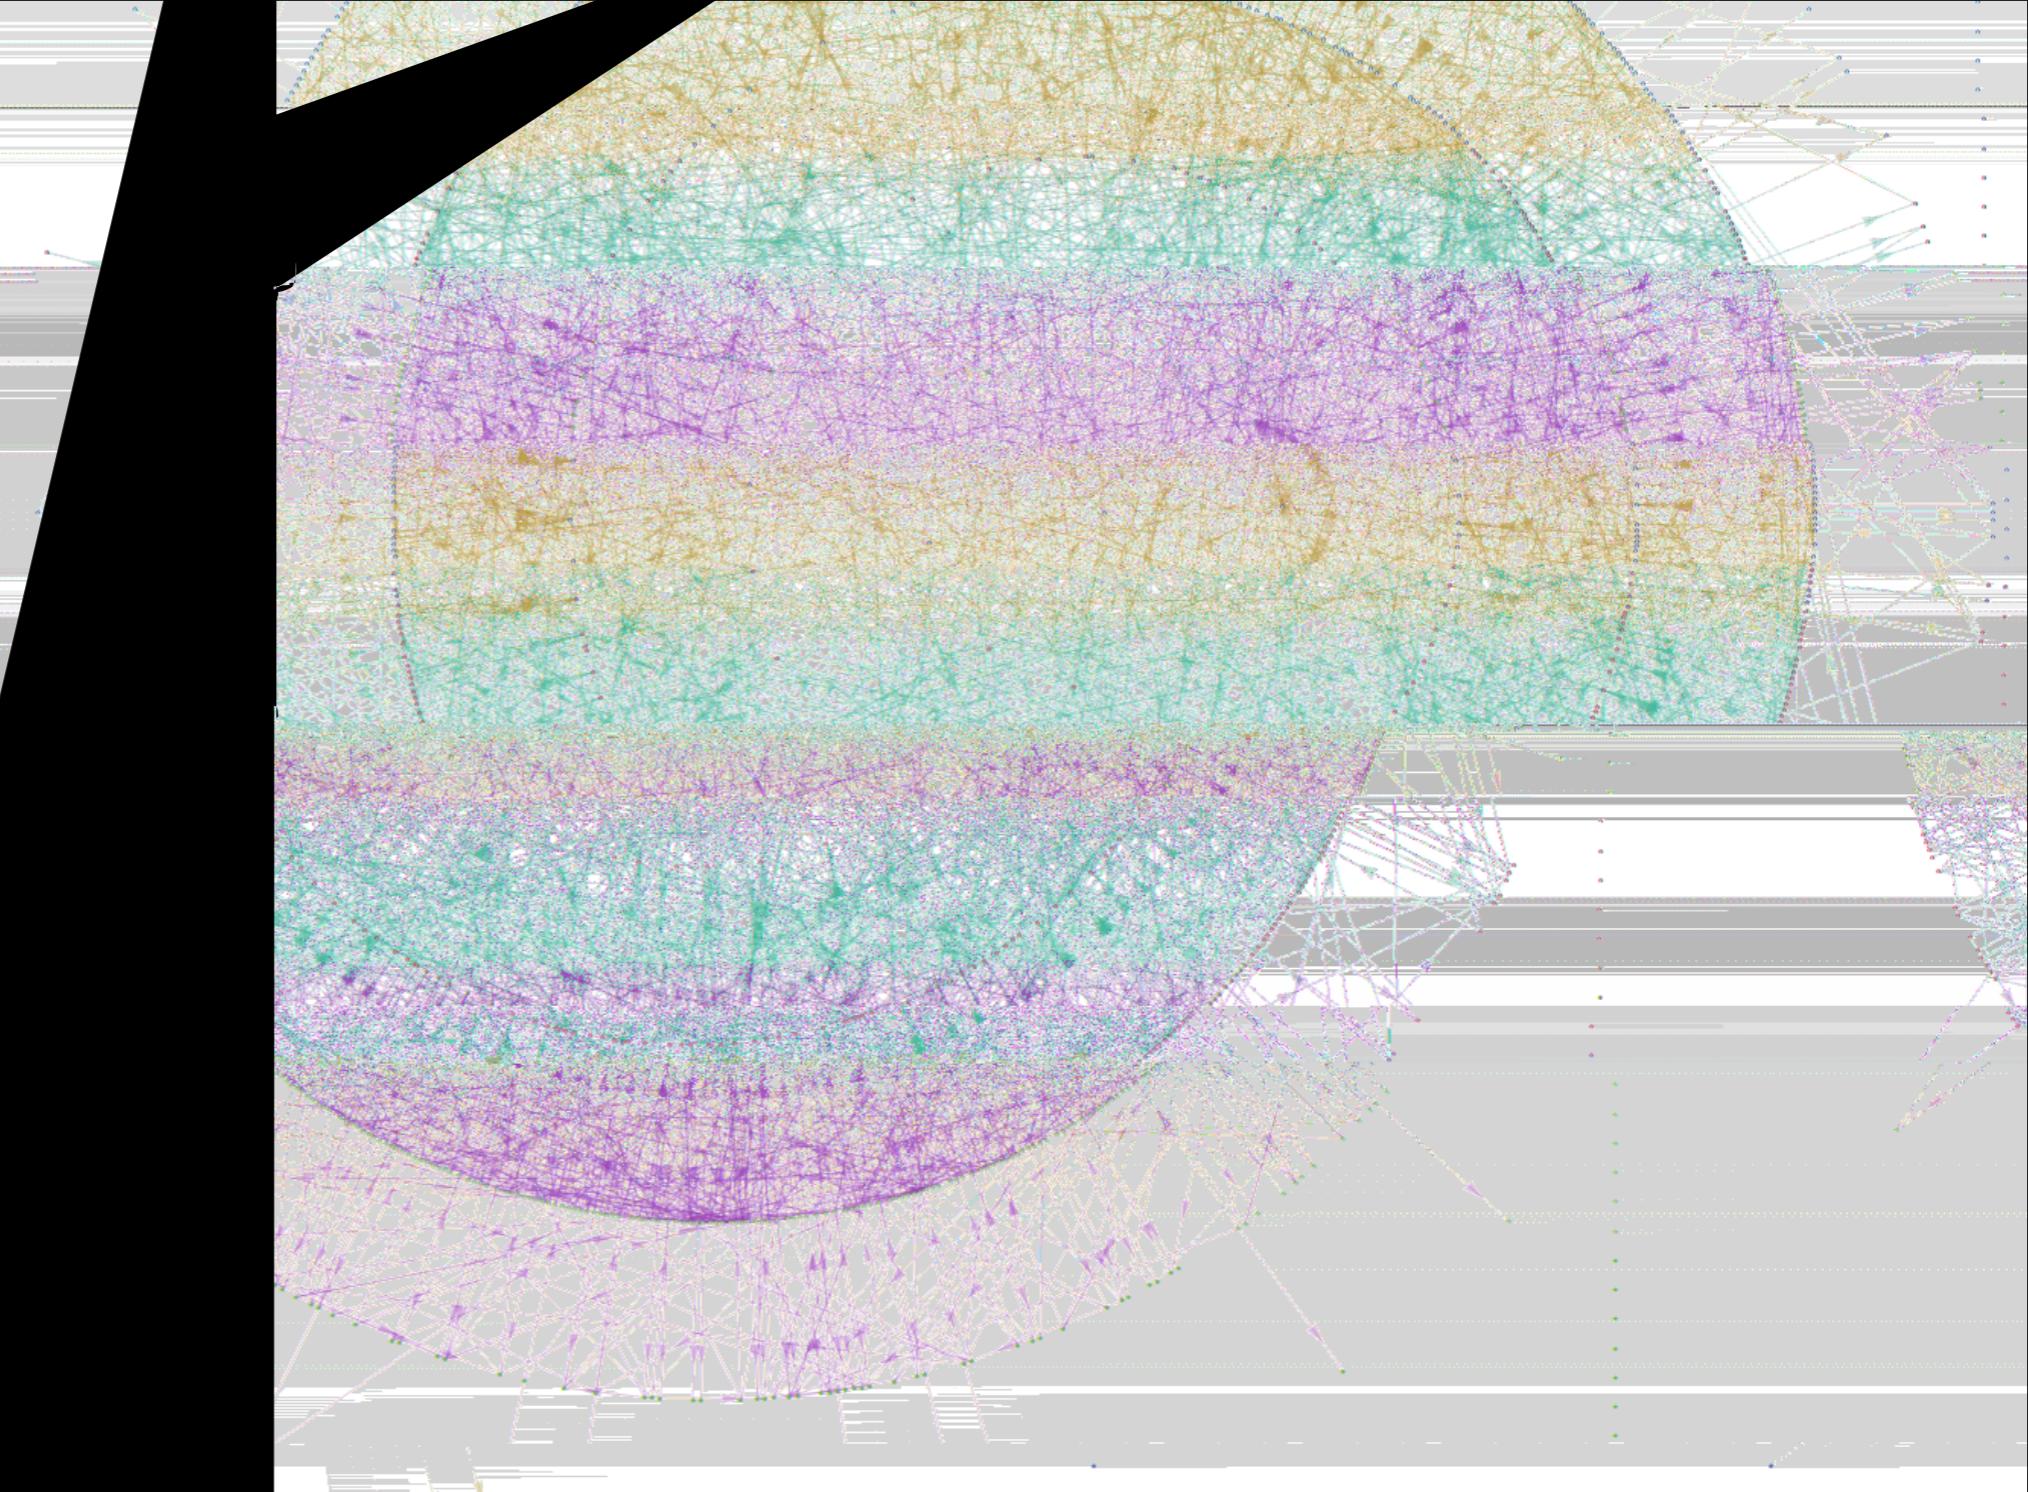


❖ Seeeed
giveeeei
feedback

❖ Seeeee
feedback b d
giveaa

❖ Seeeee
feedback b
eceiveaa

Peer Learning



Relies on Participation

Data Analysis Opportunities

SGEMM Implementation

```

#define BLOCK_SIZE 16

__global__ void MatMulKernel(float *DA, float *DB, float *DC, int Ah, int Aw,
                            int AwTiles, int Bh, int Bw) {
    // Block row and column
    int blockRow = blockIdx.y;
    int blockCol = blockIdx.x;

    // Thread row and column within Csub
    int row = threadIdx.y;
    int col = threadIdx.x;

    int cellRow = blockRow * BLOCK_SIZE + row;
    int cellCol = blockCol * BLOCK_SIZE + col;

    // Each thread computes one element of Csub
    // by accumulating results into Cvalue
    float Cvalue = 0.0;

    // Loop over all the sub-matrices of A and B that are
    // required to compute Csub
    // Multiply each pair of sub-matrices together
    // and accumulate the results

    #pragma unroll
    for (int m = 0; m < AwTiles; ++m) {

        // Shared memory used to store Asub and Bsub respectively
        __shared__ float As[BLOCK_SIZE][BLOCK_SIZE];
        __shared__ float Bs[BLOCK_SIZE][BLOCK_SIZE];

        // Load Asub and Bsub from device memory to shared memory
        // Each thread loads one element of each sub-matrix
        int aRow = BLOCK_SIZE * blockRow + row;
        int aCol = BLOCK_SIZE * m + col;

        float aValue = ((aRow < Ah) && (aCol < Aw));
        aValue *= DA[Aw * aRow + aCol];
        As[row][col] = aValue;

        int bRow = BLOCK_SIZE * m + row;
        int bCol = BLOCK_SIZE * blockCol + col;

        float bValue = ((bRow < Bh) && (bCol < Bw));
        bValue *= DB[Bw * bRow + bCol];
        Bs[row][col] = bValue;

        // Synchronize to make sure the sub-matrices are loaded
        // before starting the computation
        __syncthreads();

        // Multiply Asub and Bsub together
        for (int e = 0; e < BLOCK_SIZE; ++e)
            Cvalue += As[row][e] * Bs[e][col];

        // Synchronize to make sure that the preceding
        // computation is done before loading two new
        // sub-matrices of A and B in the next iteration
        __syncthreads();
    }

    if (cellRow >= Ah)
        return;
    if (cellCol >= Bw)
        return;

    DC[cellRow * Bw + cellCol] = Cvalue;
}

__global__ void nop() {}

// CITE: Vasily Volkov, UC Berkeley.... Interesting approach to using
// caching a bit differently... had to give it a try...
// Device SAXPY kernel
__device__ void saxpy(float a, float *b, float *c) {
    c[0] += a * b[0];
    c[1] += a * b[1];
    c[2] += a * b[2];
    c[3] += a * b[3];
    c[4] += a * b[4];
    c[5] += a * b[5];
    c[6] += a * b[6];
    c[7] += a * b[7];
    c[8] += a * b[8];
    c[9] += a * b[9];
    c[10] += a * b[10];
    c[11] += a * b[11];
    c[12] += a * b[12];
    c[13] += a * b[13];
    c[14] += a * b[14];
    c[15] += a * b[15];
}

__global__ void optimisedDLA(const float *A, int lda, const float *B, int ldb,
                            float *C, int ldc, int k) {
    const int inx = threadIdx.x;
    const int iny = threadIdx.y;
    const int ibx = blockIdx.x * 64;
    const int iby = blockIdx.y * 16;
    const int id = inx + iny * 16;

    A += ibx + id;
    B += inx + __mul24(iby + iny, ldb);
    C += ibx + id + __mul24(iby, ldc);

    const float *Blast = B + k;

    float c[16] = {0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0};

    do {
        float a[4] = {A[0 * lda], A[1 * lda], A[2 * lda], A[3 * lda]};

        __shared__ float bs[16][17];
        bs[inx][iny] = B[0 * ldb];
        bs[inx][iny + 4] = B[4 * ldb];
        bs[inx][iny + 8] = B[8 * ldb];
        bs[inx][iny + 12] = B[12 * ldb];
        __syncthreads();

        A += 4 * lda;
        saxpy(a[0], &bs[0][0], c);
        a[0] = A[0 * lda];
        saxpy(a[1], &bs[1][0], c);
        a[1] = A[1 * lda];
        saxpy(a[2], &bs[2][0], c);
        a[2] = A[2 * lda];
        saxpy(a[3], &bs[3][0], c);
        a[3] = A[3 * lda];

        A += 4 * lda;
        saxpy(a[0], &bs[4][0], c);
        a[0] = A[0 * lda];
        saxpy(a[1], &bs[5][0], c);
        a[1] = A[1 * lda];
        saxpy(a[2], &bs[6][0], c);
        a[2] = A[2 * lda];
        saxpy(a[3], &bs[7][0], c);
        a[3] = A[3 * lda];

        A += 4 * lda;
        saxpy(a[0], &bs[8][0], c);
        a[0] = A[0 * lda];
        saxpy(a[1], &bs[9][0], c);
        a[1] = A[1 * lda];
        saxpy(a[2], &bs[10][0], c);
        a[2] = A[2 * lda];
        saxpy(a[3], &bs[11][0], c);
        a[3] = A[3 * lda];

        A += 4 * lda;
        saxpy(a[0], &bs[12][0], c);
        saxpy(a[1], &bs[13][0], c);
        saxpy(a[2], &bs[14][0], c);
        saxpy(a[3], &bs[15][0], c);

        B += 16;
        __syncthreads();
    } while (B < Blast);

    for (int i = 0; i < 16; i++, C += ldc)
        C[0] = c[i];
    ...
}

```

Analysis Opportunities

- ❖ Wha e a e c ?
- ❖ Ca e de ec h e a d gi e feedbac
- ❖ Wha i i a i a e be e cia ?
- ❖ Which did e ha e he b e i
- ❖ Wha ca e GPU c a h?
- ❖ Ca e a id h e?
- ❖ Did e e agia i e?
- ❖ Wha d e he ee e ie e ?

Source of Data

- ❖ This is a big data set
- ❖ Different approaches include being efficient and different approaches
- ❖ Python packages include numpy, pandas, and tensorflow
- ❖ Python packages include numpy, pandas, and tensorflow
- ❖ NLP

I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.
I will learn these lessons for the next heterogeneous course.



What would be different?

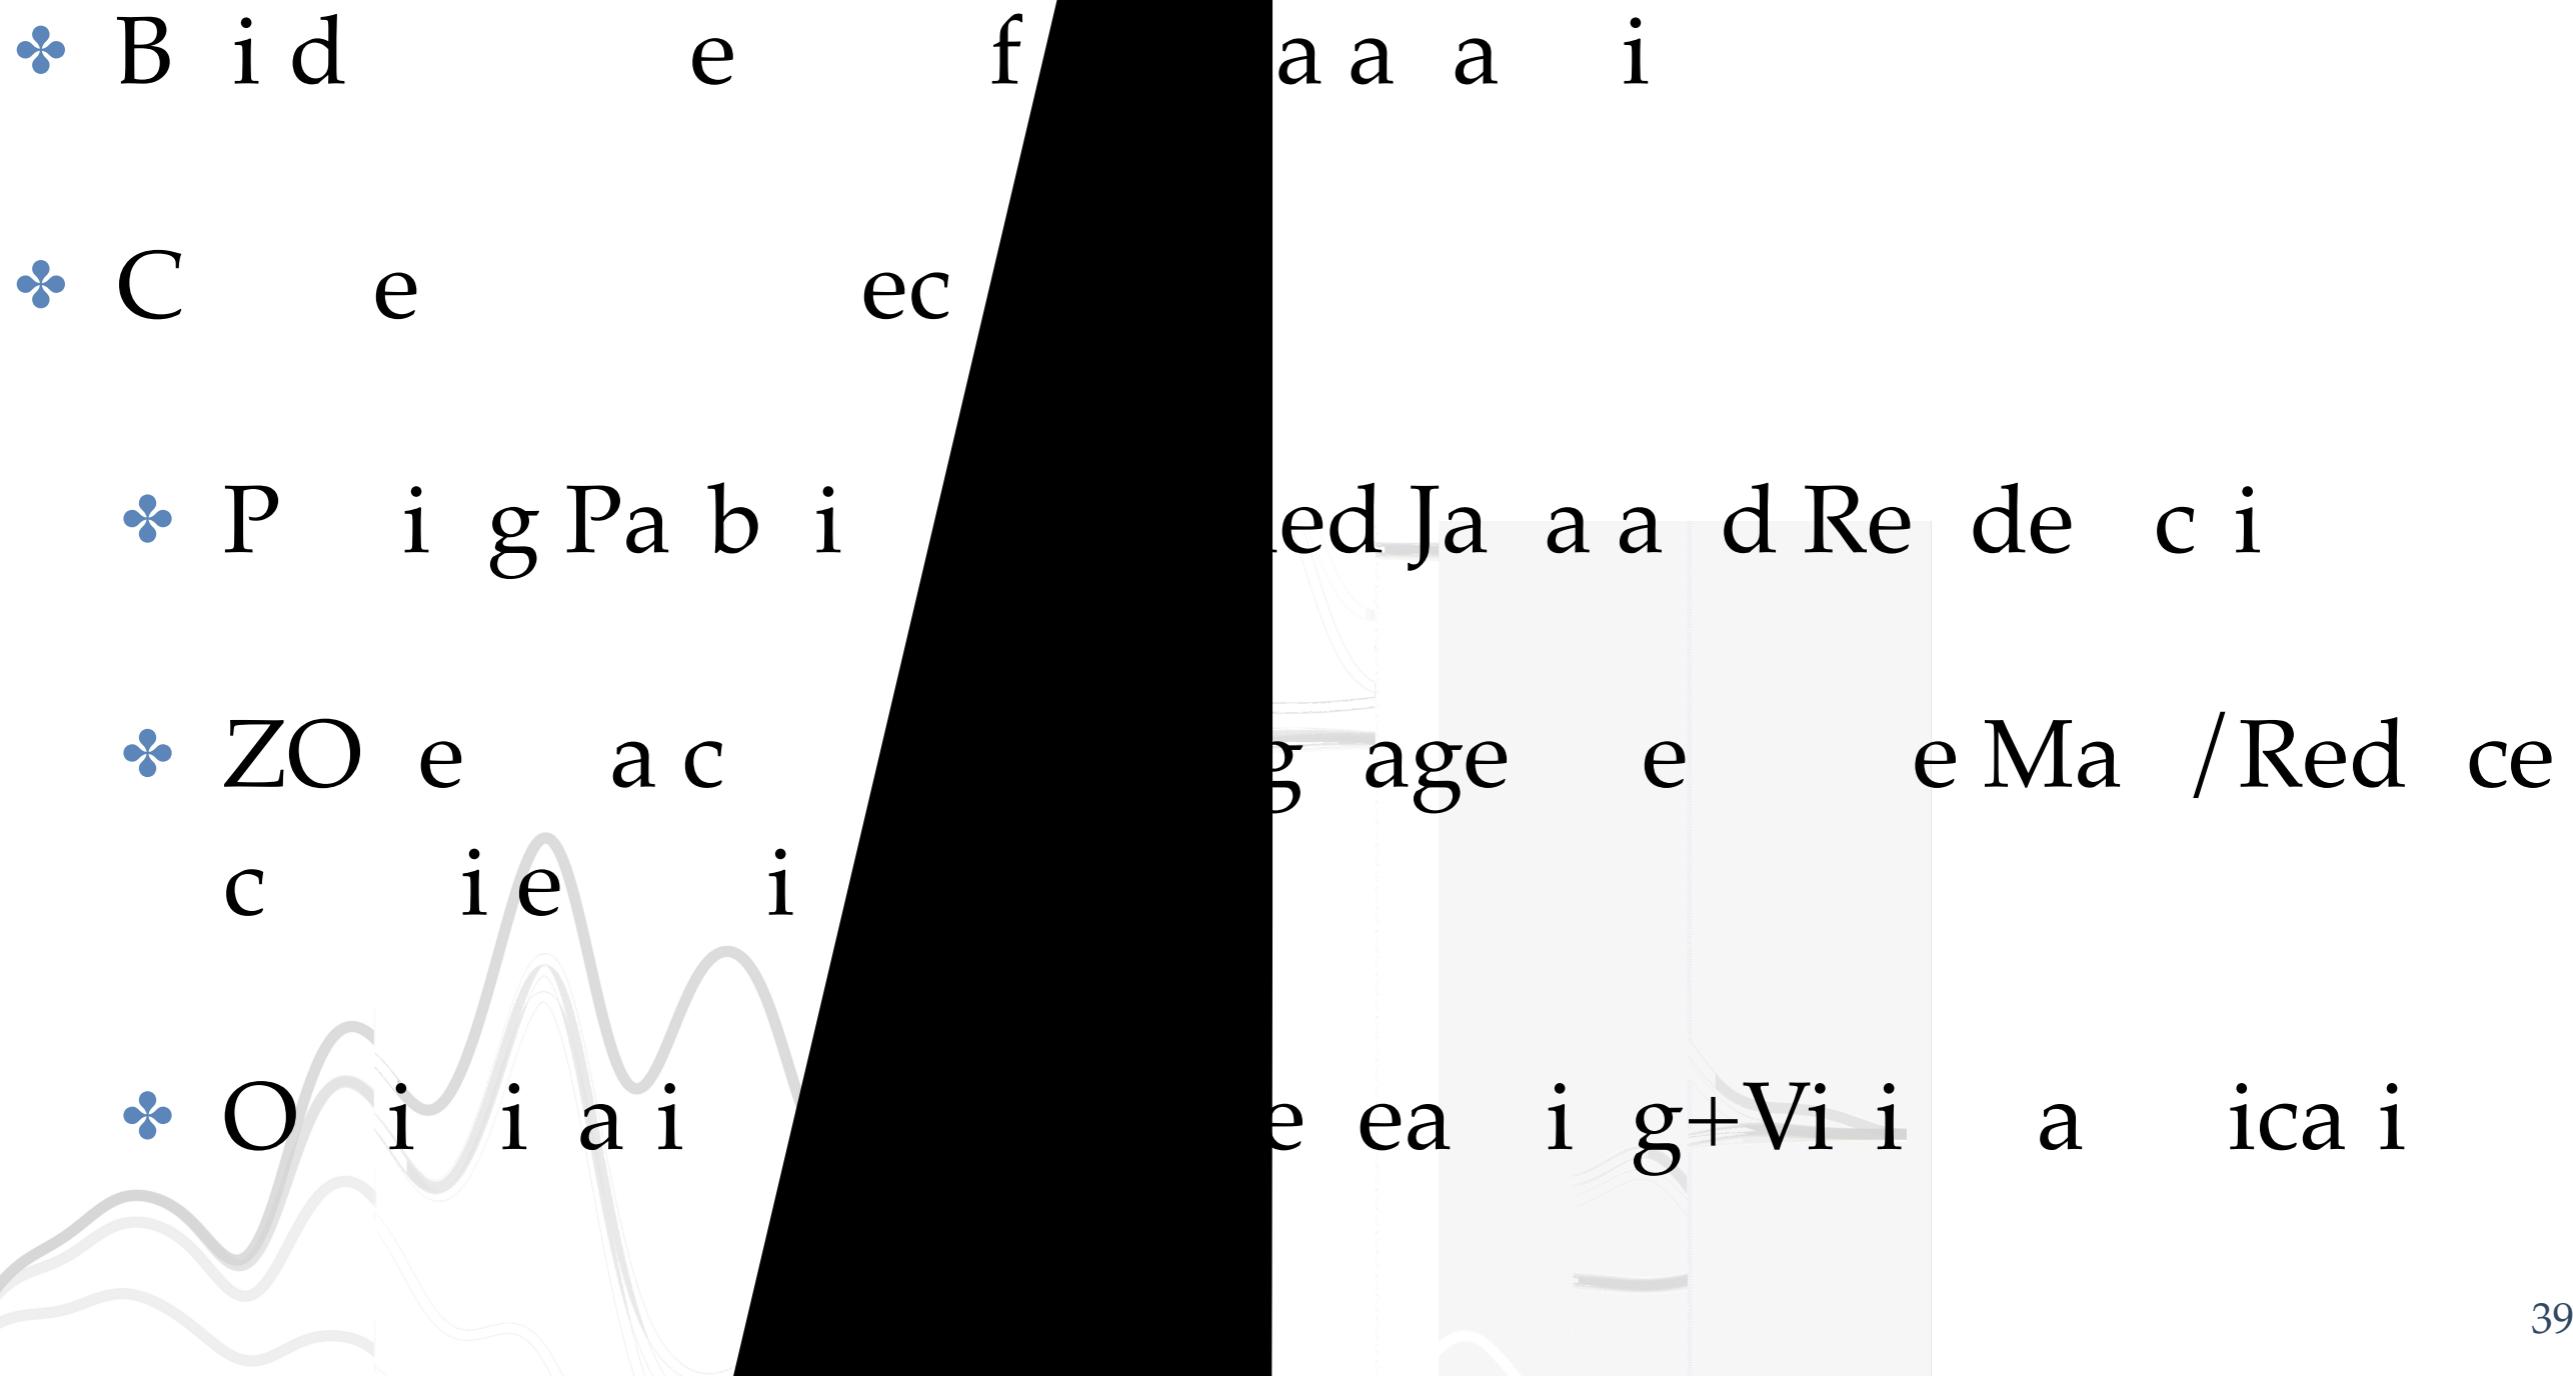
Lessons Learned

- ❖ V ee TA e e e he f a d a ed
e d e ha 30 h a da he f
- ❖ N e e e i e f e ie , 1 ea
ha e e e i ge fe ac
- ❖ S e e e c i i ci e a e f i eed
de e a hic e he



Current Work

Current Work





Questions?